

Announcement

- Lab 9 (graded)
 - To be returned this afternoon
- Assignment 6
 - Due November 11
 - Sorting: Combining C and IA-32 assembly

Lecture 25

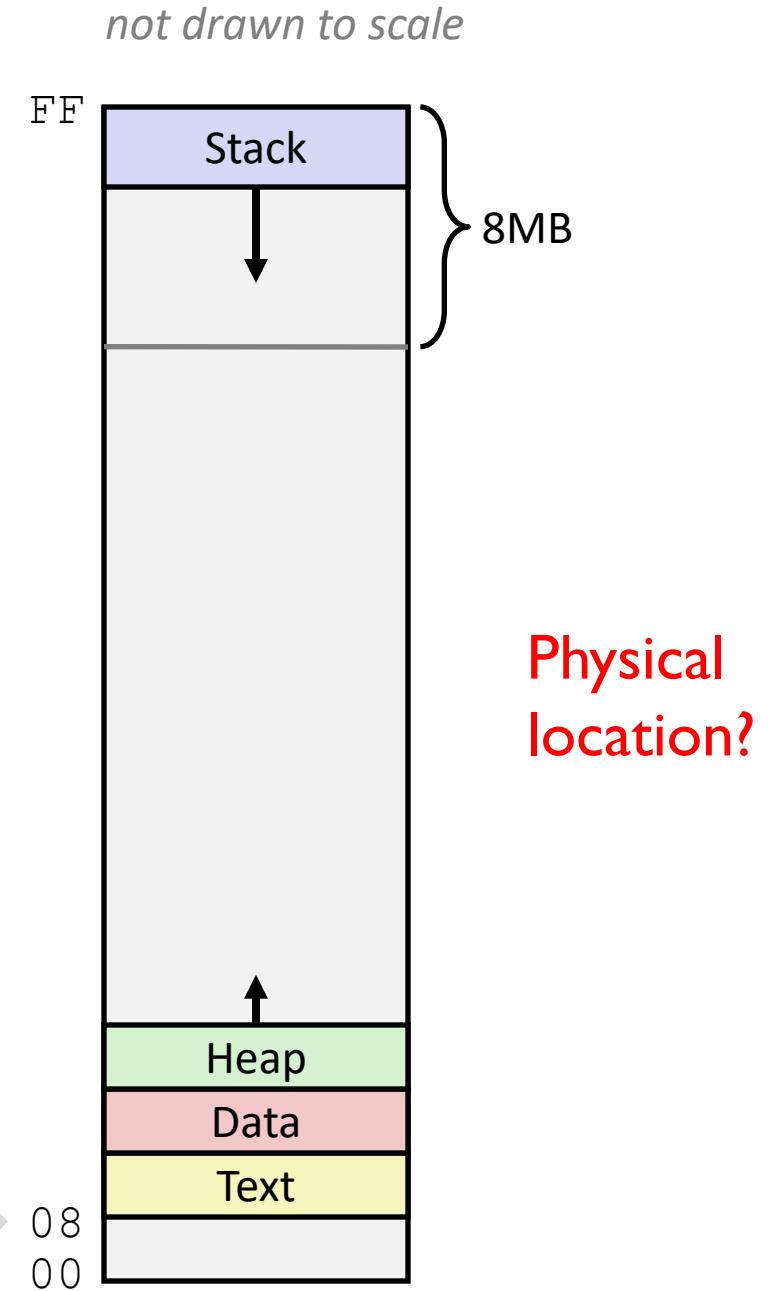
Locality

CPSC 275
Introduction to Computer Systems

IA-32 Linux Memory Layout

- Stack
 - Runtime stack (8MB limit)
 - e. g., local variables
- Heap
 - Dynamically allocated storage
 - When call malloc, calloc, new
- Data
 - Statically allocated data
 - e.g., globals & strings declared in code
- Text
 - Executable machine instructions
 - Read-only (*re-entrant*)

Upper 2 hex digits
= 8 bits of address



Random-Access Memory (RAM)

- Key features
 - RAM is traditionally packaged as a chip.
 - Basic storage unit is normally a cell (one bit per cell).
 - Multiple RAM chips form a memory.
- RAM comes in two varieties:
 - DRAM (dynamic RAM)
 - SRAM (static RAM)



SRAM vs DRAM

	Trans. per bit	Access time	Needs refresh?	Needs EDC?	Cost	Applications
SRAM	4 or 6	IX	No	Maybe	100x	Cache memories
DRAM	I	10X	Yes	Yes	IX	Main memories, frame buffers

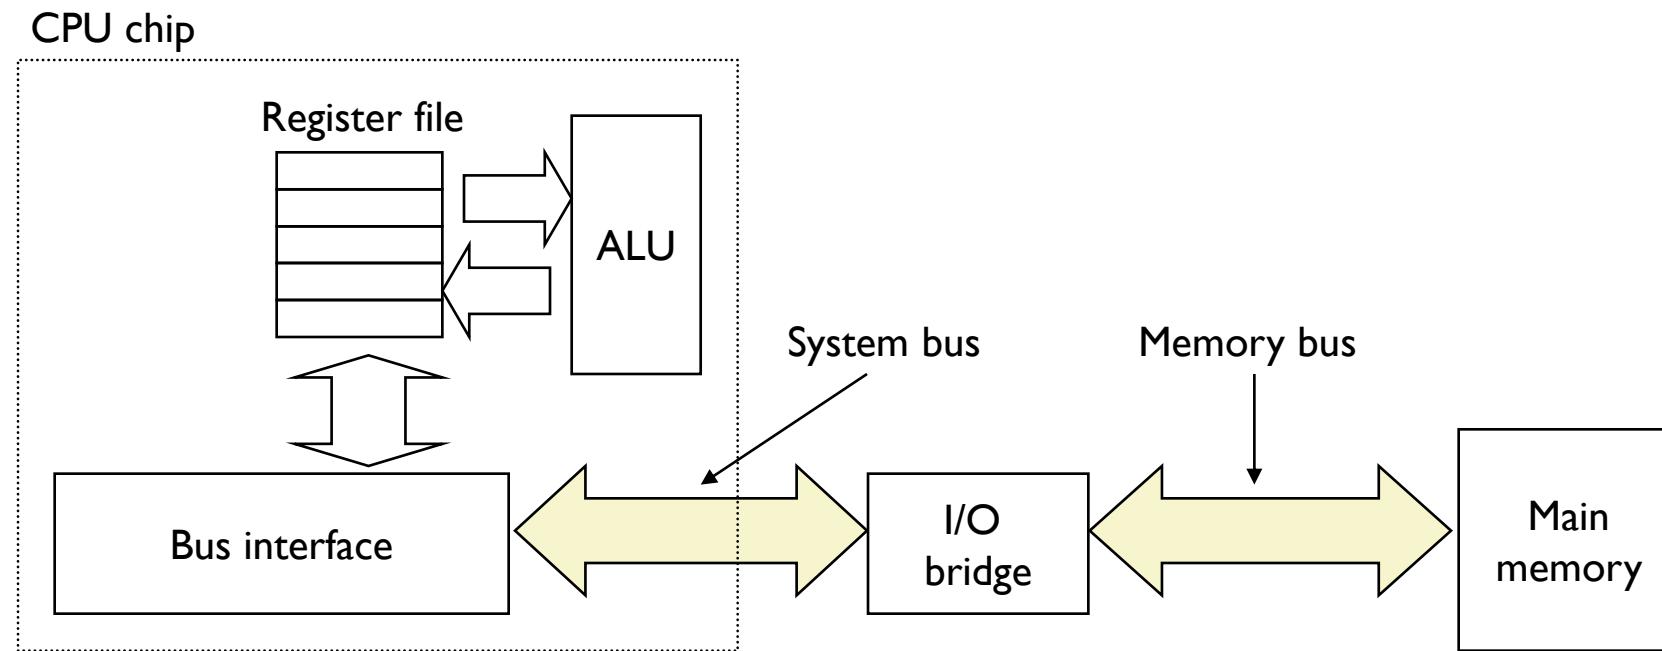
DRAM and SRAM are *volatile* memories
(lose information if powered off)

Nonvolatile Memories

- *Nonvolatile* memories retain value even if powered off
 - Read-only memory (**ROM**): programmed during production
 - Programmable ROM (**PROM**): can be programmed once
 - Eraseable PROM (**EPROM**): can be bulk erased (UV, X-Ray)
 - Electrically eraseable PROM (**EEPROM**): electronic erase capability
 - Flash memory: EEPROMs. with partial (block-level) erase capability
- Uses for nonvolatile memories
 - *Firmware* programs stored in a ROM (BIOS, controllers for disks, network cards, graphics accelerators, security subsystems,...)
 - Solid state disks (SSD)
 - Disk caches


Connecting CPU and Memory

- A **bus** is a collection of parallel wires that carry address, data, and control signals.
- Buses are typically shared by multiple devices.



The CPU-Memory Gap

The gap widens between DRAM, disk, and CPU speeds.


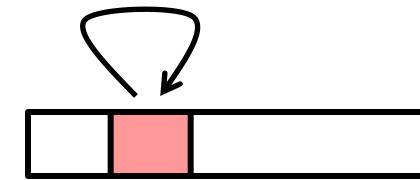


Locality to the rescue!

The key to bridging this CPU-Memory gap is a fundamental property of computer programs known as **locality**.

What is Locality?

- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently
- **Temporal locality:**
 - Recently referenced items are likely to be referenced again in the near future
- **Spatial locality:**
 - Items with nearby addresses tend to be referenced close together in time



Locality Example

```
sum = 0;  
for (i = 0; i < n; i++)  
    sum += a[i];  
return sum;
```

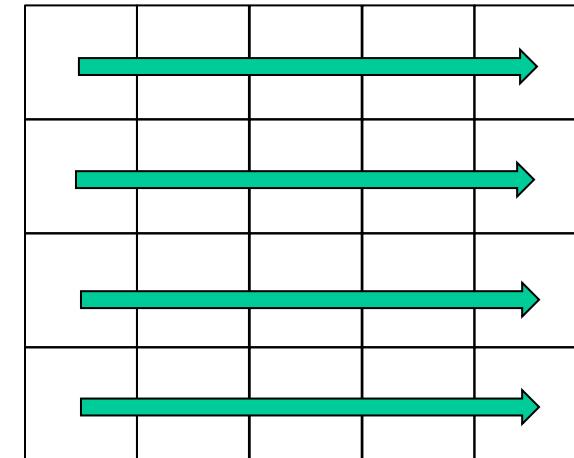
- data references
 - Reference array elements in succession (stride-1 reference pattern). spatial locality
 - Reference variable **sum** each iteration. temporal locality
- instruction references
 - Reference instructions in sequence. spatial locality
 - Cycle through loop repeatedly. temporal locality

Qualitative Estimates of Locality

- Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer.
- **Q:** Does this function exhibit good locality with respect to array *a*?

```
int sum_array_rows(int a[M] [N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i] [j];
    return sum;
}
```




Q: Stride? 1

Locality Example

- **Q:** Does this function exhibit good locality with respect to array a?

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```



Q: Stride? N

- **Q:** Which one is better?

Locality Example

- **Q:** Can you permute the loops so that the function scans the 3-d array `a` with a stride-1 reference pattern and thus has good spatial locality? (Homework)


```
int sum_array_3d(int a[N] [N] [N] )
{
    int i, j, k, sum = 0;

    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            for (k = 0; k < N; k++)
                sum += a[k] [i] [j];
    return sum;
}
```

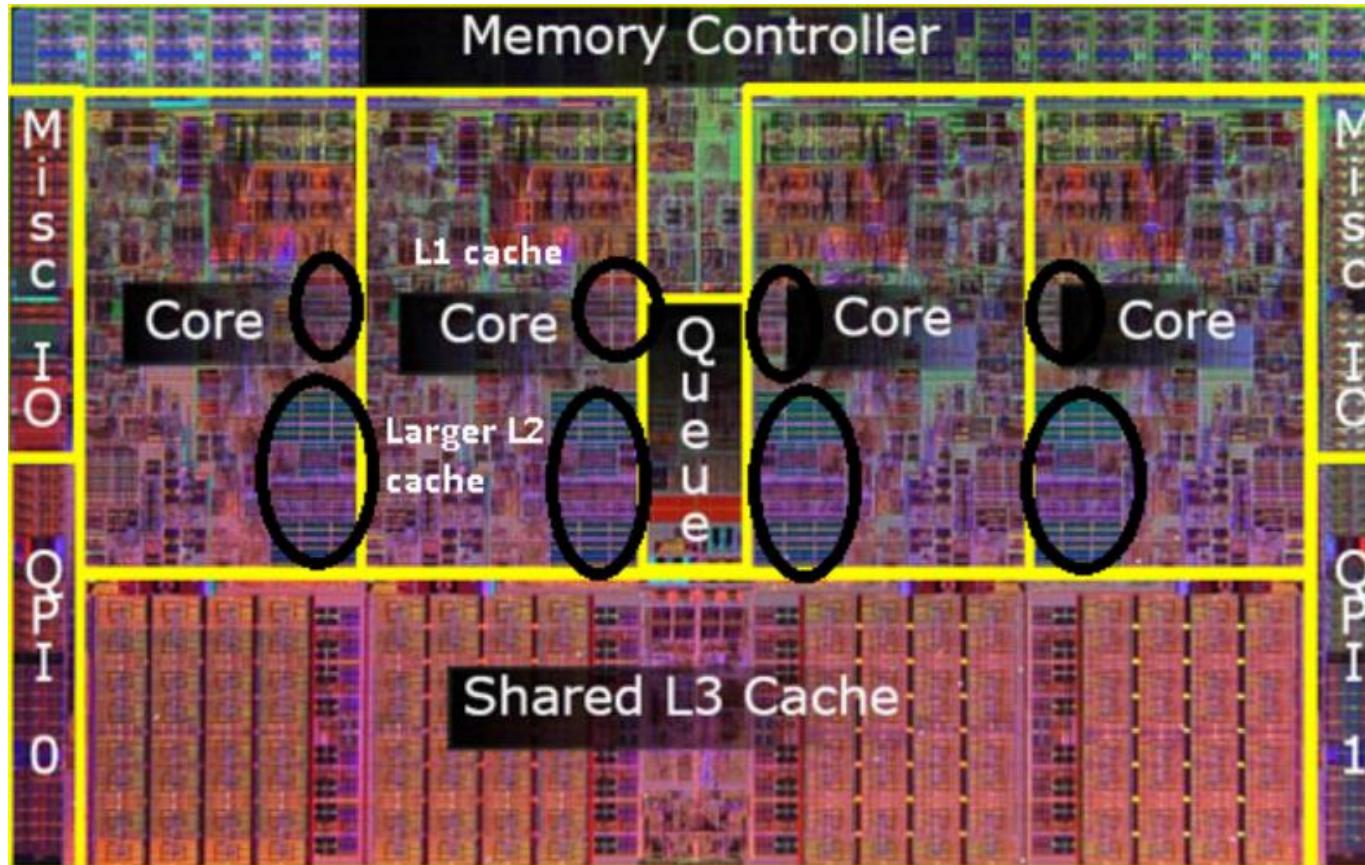
Memory Hierarchy

- Some fundamental and enduring properties of hardware and software:
 - Fast storage technologies cost more per byte, have less capacity, and require more power (heat!).
 - The gap between CPU and main memory speed is widening.
 - **Well-written** programs tend to exhibit **good locality**.
- This led to an approach of organizing memory and storage systems known as a **memory hierarchy**.

Memory Hierarchy



Where is cache?



A typical quad-core processor (Intel i7)

Cache

- **Cache:** A smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device.
- Fundamental idea of a memory hierarchy:
 - For each k , the faster, smaller device at level k serves as a cache for the larger, slower device at level $k+1$.
- Why do memory hierarchies work?
 - Because of locality, programs tend to access the data at level k more often than they access the data at level $k+1$.
 - Thus, the storage at level $k+1$ can be slower, and thus larger and cheaper per bit.

