

# Announcement


- Assignment 7
  - Due November 18
  - Cache simulator

Lecture 27


# More on Cache

CPSC 275  
Introduction to Computer Systems

# General Cache Organization ( $S, E, B$ )




# Cache Read



# Example: Direct Mapped Cache ( $E = 1$ )

Direct mapped: One line per set


Assume: cache block size 8 bytes ( $b = ?$ )



# Example: Direct Mapped Cache ( $E = 1$ )

Direct mapped: One line per set


Assume: cache block size 8 bytes



# Example: Direct Mapped Cache ( $E = 1$ )

Direct mapped: One line per set

Assume: cache block size 8 bytes



If no match, an old line is evicted and replaced.

# Direct-Mapped Cache Simulation

4-bit address

x	xx	x
---	----	---

t=1    s=2    b=1

M=16 byte addresses (total size of memory)

B=2 bytes/block

S=4 sets

E=1 line/set

# Direct-Mapped Cache Simulation

4-bit address

x	xx	x
---	----	---

t=1    s=2    b=1

M=16 byte addresses (total size of memory)

B=2 bytes/block

S=4 sets

E=1 line/set

Address trace:

0	[0000 <sub>2</sub> ]	miss
1	[0001 <sub>2</sub> ]	hit
7	[0111 <sub>2</sub> ]	miss
8	[1000 <sub>2</sub> ]	miss
0	[0000 <sub>2</sub> ]	miss


	v	Tag	Block
Set 0	1	0	M[0-1]
Set 1			
Set 2			
Set 3	1	0	M[6-7]

# 2-way Set Associative Cache


$E = 2$ : Two lines per set

Assume: cache block size 8 bytes


Address of short int:



find set




• •



# 2-way Set Associative Cache

$E = 2$ : Two lines per set


Assume: cache block size 8 bytes



# 2-way Set Associative Cache

E = 2: Two lines per set


Assume: cache block size 8 bytes



No match:

- One line in set is selected for eviction and replacement
- Replacement policies: random, least recently used (LRU), ...

# 2-Way Set Associative Cache Simulation




M=16 byte addresses,  
B=2 bytes/block,  
S=2 sets,  
E=2 lines/set

Address trace (reads, one byte per read):

0	[0000 <sub>2</sub> ]	miss
1	[0001 <sub>2</sub> ]	hit
7	[0111 <sub>2</sub> ]	miss
8	[1000 <sub>2</sub> ]	miss
0	[0000 <sub>2</sub> ]	hit

	v	Tag	Block
Set 0	1	00	M[0-1]
	1	10	M[8-9]
Set 1	1	01	M[6-7]
	0		

# Fully Associative Cache ( $S = 1$ )





# Assignment 7: Cache Simulator

- A C program that simulates the behavior of a cache memory
  - Parameters:  $s$ ,  $E$ ,  $b$
  - Given a memory access *trace file* as input,
    - Simulates the hit/miss behavior of a cache memory on this trace and
    - Prints the total number of hits, misses, and evictions.

# Trace Files

- Trace files generated by valgrind, a Linux utility program for memory debugging, memory leak detection, and profiling
  - Example trace file

I 0400d7d4, 8	// instruction load
M 0421c7f0, 4	// modify (a load followed by a store)
L 04f6b868, 8	// load
S 7ff0005c8, 8	// store

↑      ↑      ↑  
space    address    data size

# Trace Files, cont'd

- For this assignment, we are interested only in data cache performance, so your simulator should ignore all instruction cache accesses (lines starting with “I”).
- To help you simplify the code, all lines with an instruction load have been already removed from trace files.

# LRU Replacement Policy

- Necessary conditions for replacement
  - a cache miss occurred
  - the current set is full
- Need to determine which line from the current set will be replaced by a new line (*eviction*).
- Pick the one *least recently used* (referenced).
  - Why does it work?
  - How to implement?

# Running Your Simulator

```
$ ./mycache -s 2 -E 1 -b 4 -t traces/tiny.trace
```

To verify the result:

```
$ ./refcache -s 2 -E 1 -b 4 -t traces/tiny.trace
```

# Programming Notes

- For arbitrary  $s$ ,  $E$ , and  $b$ , you will need to allocate storage for your simulator's data structures using `malloc`.
- 64-bit address field: `unsigned long int`
- Ignore the request data sizes in the traces, that is, there is no need to define blocks in lines.
- Use a timestamp to keep track of the last time each line has been referenced.
- Start with small trace files.



# Command-Line Arguments

- When we run a program, we'll often need to supply it with information.
- Example:

```
$ ./repeat 10 computer
```

# Command-Line Arguments

- To obtain access to *command-line arguments*, main must have two parameters:

```
int main(int argc, char *argv[] )  
{  
    ...  
}
```

# Command-Line Arguments


- `argc` (“argument count”) is the number of command-line arguments.
- `argv` (“argument vector”) is an array of pointers to the command-line arguments (stored as strings).
- `argv[0]` points to the name of the program,
- `argv[1]` through `argv[argc-1]` point to the remaining command-line arguments.

# Command-Line Arguments

- If the user enters the command line

```
$ ./repeat 10 computer
```

then argc will be 3 (why?),  
and argv will have the following appearance:



# Command-Line Arguments

- Since `argv` is an array of pointers, accessing command-line arguments is easy, e.g,

```
int i;  
for (i = 1; i < argc; i++)  
    printf("%s\n", argv[i]);
```

## Output:

```
./repeat  
10  
computer
```

- For numeric arguments, conversion might be necessary:

```
int count = atoi(argv[1]);
```

# Command-Line Options

- Command-line *options* modify the program's behavior.
- Command-line arguments and options

```
$ ls -l myfile
```

*option argument*

- Some command-line options take values:

```
$ tail -n 20 myfile
```

- The **getopt** function is useful in parsing command-line options.

# The getopt Function

```
#include <unistd.h>

int getopt(int argc, char *argv[], char *optstring);
extern char *optarg;
```

- **argc** is the number of arguments.
- **argv** is an array of arguments.
- **optstring** is a string containing the option characters.
  - If such a character is followed by a colon, the option requires a value.
- **optarg** is an *external* variable string containing the option values.
- If there are no more option characters, the function returns -1.

# The getopt Function


- Normally, getopt is called in a loop.
- When getopt returns -1, indicating no more options are present, the loop terminates.
- A switch statement is used to dispatch on the return value from getopt.

# Example using getopt

```
#include <stdio.h>
#include <unistd.h>

int main (int argc, char *argv[]) {
    int r, cval;
    while ((r = getopt(argc, argv, "ab:c:")) != -1)
        switch (r) {
            case 'a':
                printf("Do option a!\n"); break;
            case 'b':
                printf("Do option b with %s!\n", optarg); break;
            case 'c':
                cval = atoi(optarg);
                printf("Do option c with %d!\n", cval);
                break;
            default:
                printf("Error: Unknown option!\n");
        }
}
```

three options: a, b, and c  
b and c require values  
c's value is a number



# Opening a file

```
#include <stdlib.h>
FILE *fopen(char *filename, char *mode);
```

- *filename* is the name of the file to be opened.
- *mode* is a “mode string” that specifies what operations we intend to perform on the file.
- returns a *file pointer* (or **NULL** on error):

```
FILE fp = fopen("data.in", "r");
... // do something with the file
fclose(fp);
```

# Modes

Mode strings for text files:

<b>String</b>	<b>Meaning</b>
"r"	Open for reading
"w"	Open for writing (file need not exist)
"a"	Open for appending (file need not exist)
"r+"	Open for reading and writing, starting at beginning
"w+"	Open for reading and writing (truncate if file exists)
"a+"	Open for reading and writing (append if file exists)

# Reading from a file

```
char *fgets(char *buf, int n, FILE *stream)
```

- *buf* is the pointer to an array of chars where the string read is stored.
- *n* is the maximum number of characters to be read, including the null character.
- *stream*, is the pointer to a FILE object
- Example:

```
FILE fp = fopen("data.in", "r");  
char buf[10];  
fgets(buf, 10, fp);
```

# Reading from a string

```
int sscanf(char *str, char *format, ...)
```

- *str* is the pointer to an array of chars where the string is stored.
- *format* is a format specifier
- Example:

```
char s[] = "November 20, 2023";
char mon[20];
int day, yr;
sscanf(s, "%s %d, %d", mon, &day, &yr);
```

→ mon = “November”, day = 20, yr = 2023

