

# Announcements

- Assignment 5
  - Due November 4
  - String operations in IA-32
- CPSC 316: Foundation of Programming Languages
  - Study of PLs: design, evaluation, and implementation
  - Compiler construction
- CPSC 375: High-Performance Computing
  - Parallel algorithms
  - Concurrent programming
  - Distributed computing
  - GPU programming

Lecture 22

# Data Alignment

CPSC 275  
Introduction to Computer Systems

# What's the Size?

- Consider

```
struct S {  
    char c;  
    int i;  
    short s;  
};
```

**Q:** What is sizeof(S) under IA-32?

**A:** 12 (??)

# What Is Data Alignment?

- Storing variables at memory addresses that are multiples of their sizes.
- Example:
  - int (4 bytes) stored at address 0x1000 (multiple of 4).
  - short (2 bytes) stored at address 0x1002 (multiple of 2).
- Unaligned access: occurs when data crosses a boundary not matching its size (e.g., int at 0x1003).

# Memory and Data in IA-32

- IA-32 operates with a 32-bit address bus, providing a linear, 4-gigabyte (4GB) address space
- The smallest unit of memory that can be individually addressed is a single byte.
- The architecture handles data in various sizes (1, 2, 4, or 8).
  - 4 bytes on IA-32
- Data is fetched from memory in fixed-size "chunks" over bus (32 bits).
  - Aligned access → 1 memory operation
  - Misaligned access → 2 or more memory operations
  - On older CPUs: could even trigger a bus error or fault.

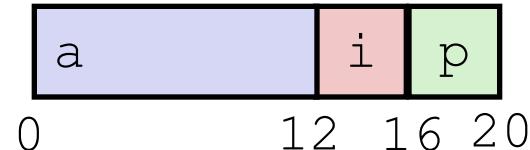
# Why Data Alignment?

- Best performance when data sizes align perfectly with the boundaries of the memory bus
  - Allows single, efficient transactions.
- Unaligned memory access allowed, but suboptimal performance
  - CPU must perform extra work to retrieve data that spans across its natural memory access boundaries
- Accessing data that crosses *cache lines* results in *cache misses* and multiple memory cycles (TBD).

# Structure Allocation

```
struct rec {  
    int a[3];  
    int i;  
    char *p;  
};
```

Memory Layout



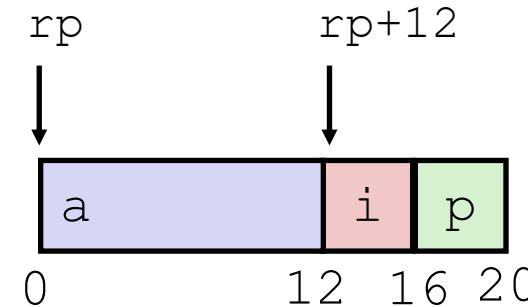
- Concept
  - Contiguously-allocated region of memory
  - Members may be of different types
  - Refer to members within structure by names
  - Examples:

```
struct rec r;  
r.i = 10;
```

```
struct rec *rp = &r;  
rp->i = 10;
```

# Structure Access

```
struct rec {  
    int a[3];  
    int i;  
    char *p;  
};
```



## ■ Accessing Structure Member

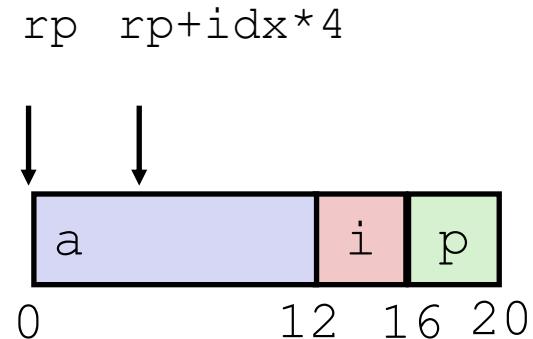
- Pointer indicates first byte of structure
- Access elements with offsets

```
void set_i(struct rec *rp, int val) {  
    rp->i = val;  
}
```

```
# %edx = val, %eax = rp  
movl %edx, 12(%eax) # Mem[rp+12] = val
```

# Pointer to Structure Member

```
struct rec {  
    int a[3];  
    int i;  
    char *p;  
};
```

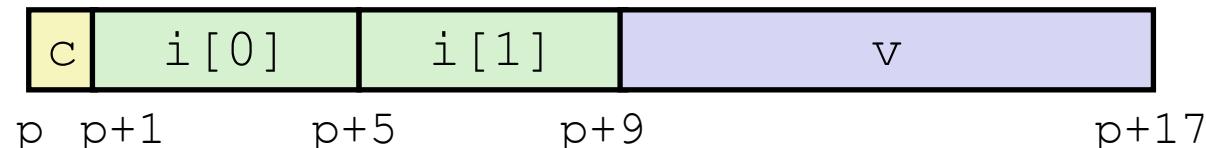


```
int *get_ap(struct rec *rp, int idx) {  
    return &(rp->a[idx]);  
}
```

```
movl 12(%ebp), %eax      # get idx  
sall $2, %eax            # idx*4  
addl 8(%ebp), %eax      # r+idx*4
```

# Structures and Alignment

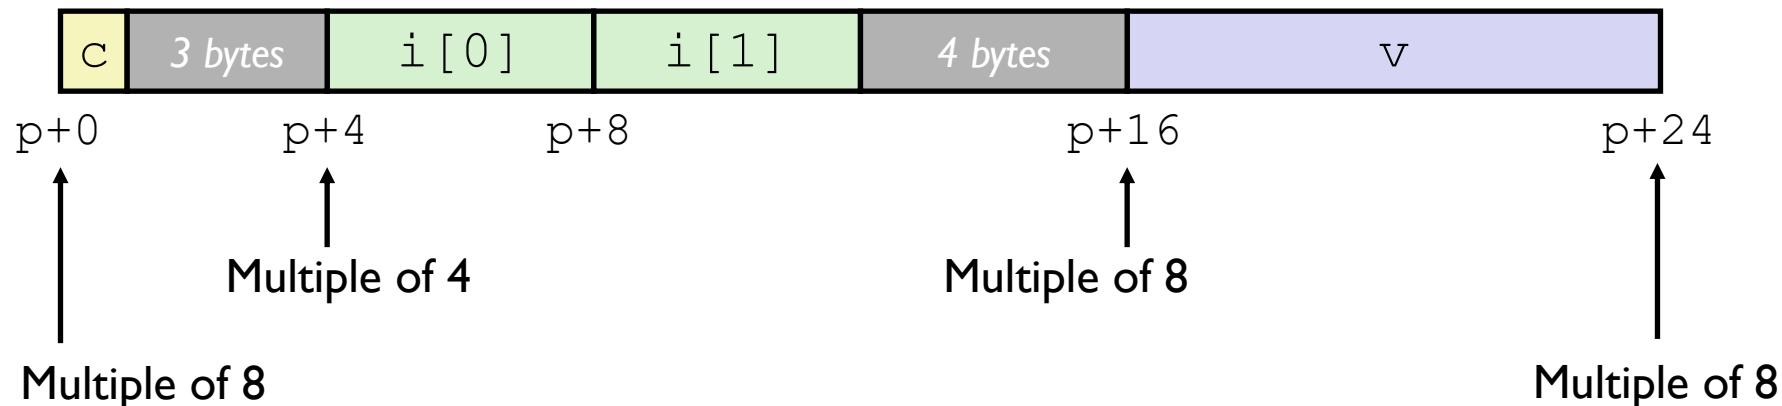
## ■ Unaligned data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

## ■ Aligned data

- Primitive data type requires  $K$  bytes
- Address must be multiple of  $K$



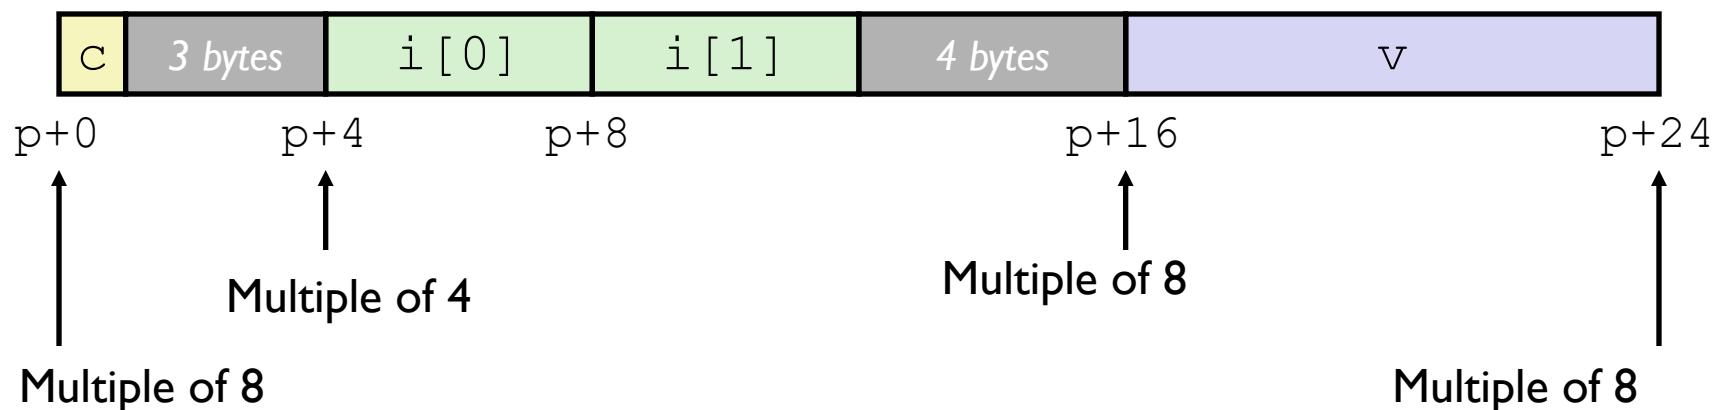
# Alignment in IA-32

- 1 byte: **char**, ...
  - no restrictions on address
- 2 bytes: **short**, ...
  - lowest 1 bit of address must be  $0_2$
- 4 bytes: **int**, **float**, **char \***, ...
  - lowest 2 bits of address must be  $00_2$
- 8 bytes: **double**, ...
  - Windows:
    - lowest 3 bits of address must be  $000_2$
  - Linux:
    - lowest 2 bits of address must be  $00_2$
    - i.e., treated the same as a 4-byte primitive data type

# Satisfying Alignment with Structures

- Within structure:
  - Must satisfy each element's alignment requirement
- Overall structure placement
  - Each structure has alignment requirement  $K$ 
    - $K = \text{largest alignment of any element}$
  - Initial address and structure length must be multiples of  $K$
- Example (Windows):
  - $K = 8$ , due to **double** element

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

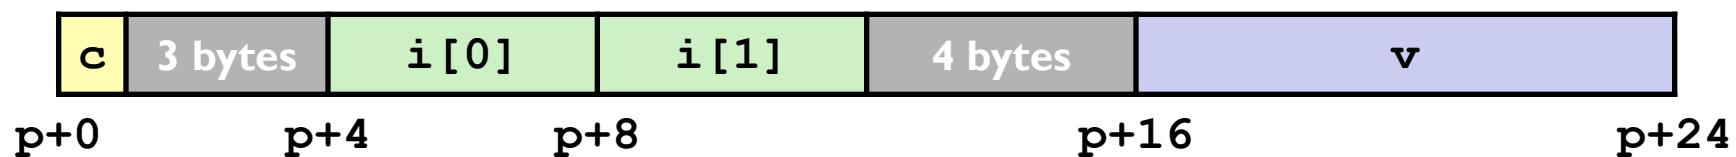


# Different Alignment Conventions

- IA-32 Windows:

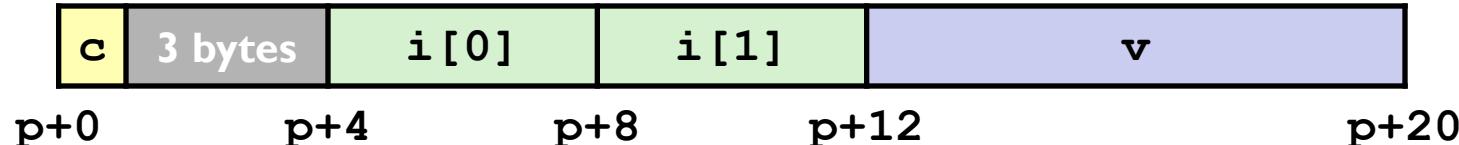
- $K = 8$ , due to **double** element

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```



- IA-32 Linux

- $K = 4$ ; **double** treated like a 4-byte data type

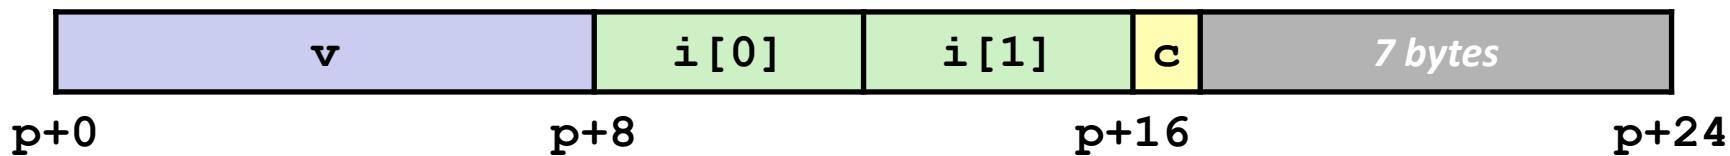


# Meeting Overall Alignment Requirement

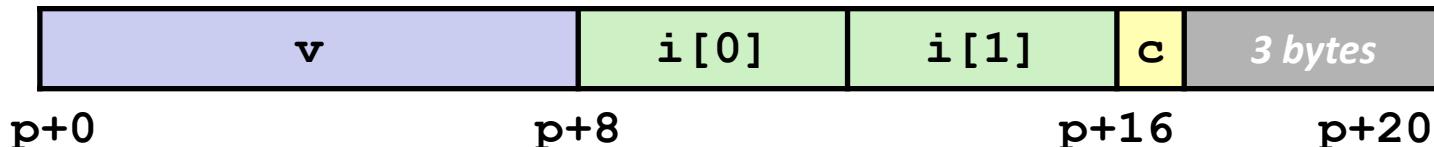
- For largest alignment requirement  $K$
- Overall structure must be multiple of  $K$

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```

- IA-32 Windows:



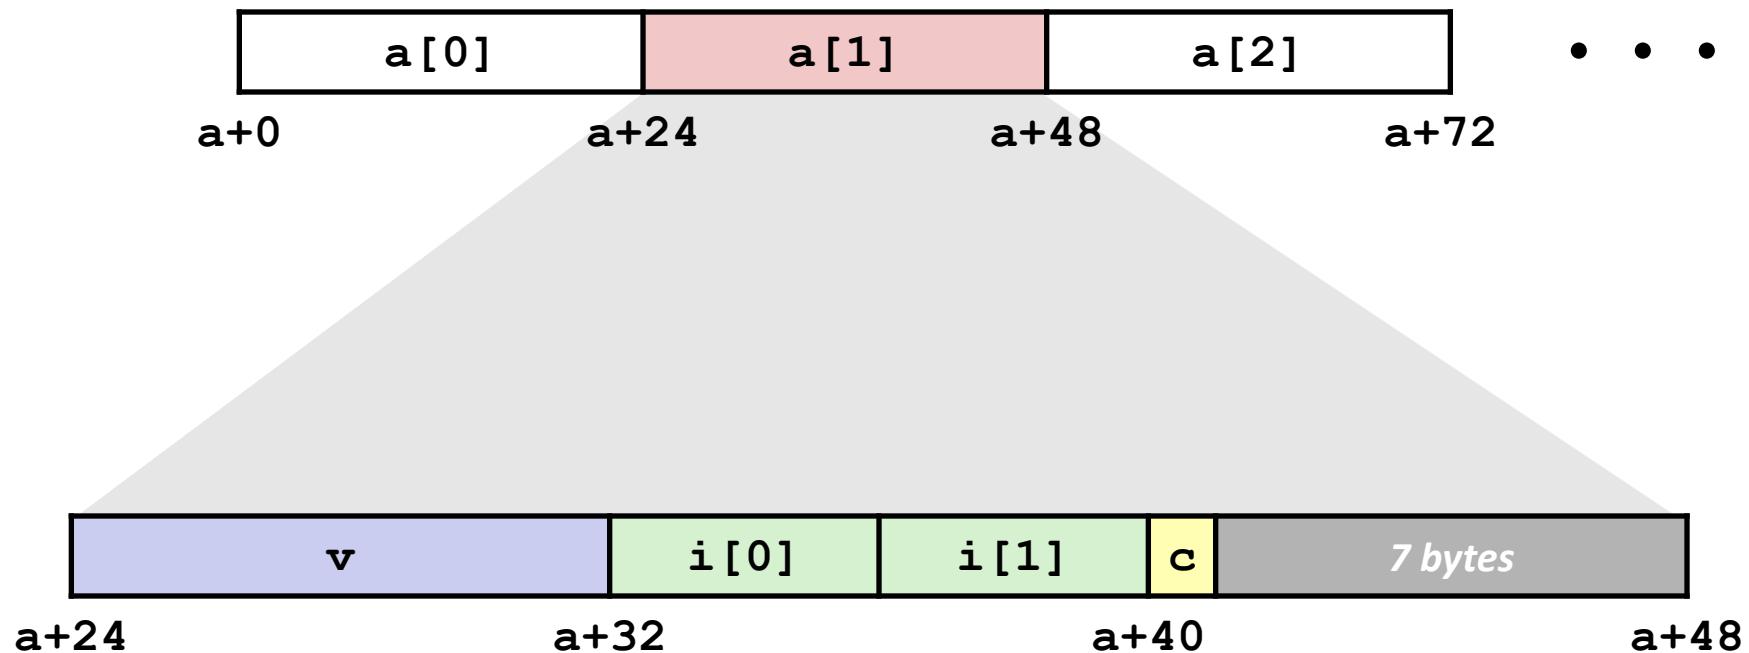
- IA-32 Linux



# Arrays of Structures

- Overall structure length multiple of  $K$
- Satisfy alignment requirement for every element
- Example: IA-32 Windows

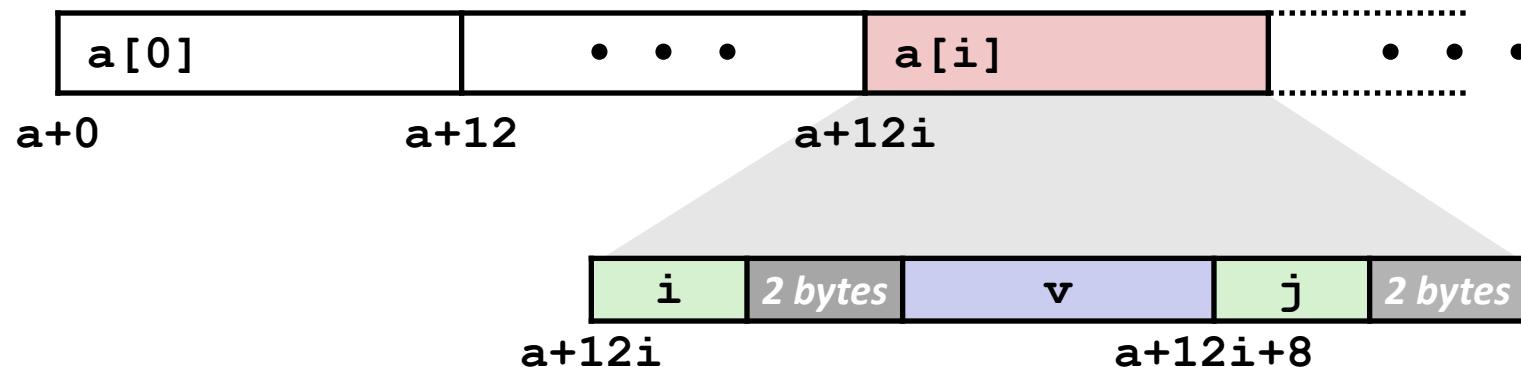
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



# Accessing Array Elements

- What is the value of `sizeof(struct S3)`?
- Element `j` is at offset 8 within structure
- Assembler gives offset `a+8`
- Compute array offset `12i` for element `i`

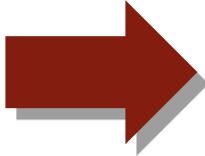
```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```



# Saving Space

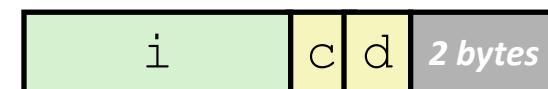
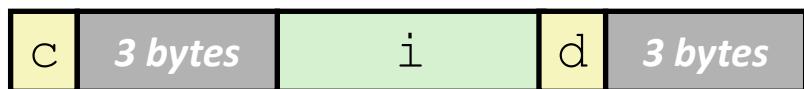
- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- Effect ( $K = 4$ )



# Union Allocation

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```



- Allocate according to largest element
- Can only use one field at a time

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```

