

Announcement

- Lab 9 (graded)
 - To be returned this afternoon
- Assignment 6
 - Due November 11
 - Sorting: Combining C and IA-32 assembly

Lecture 25

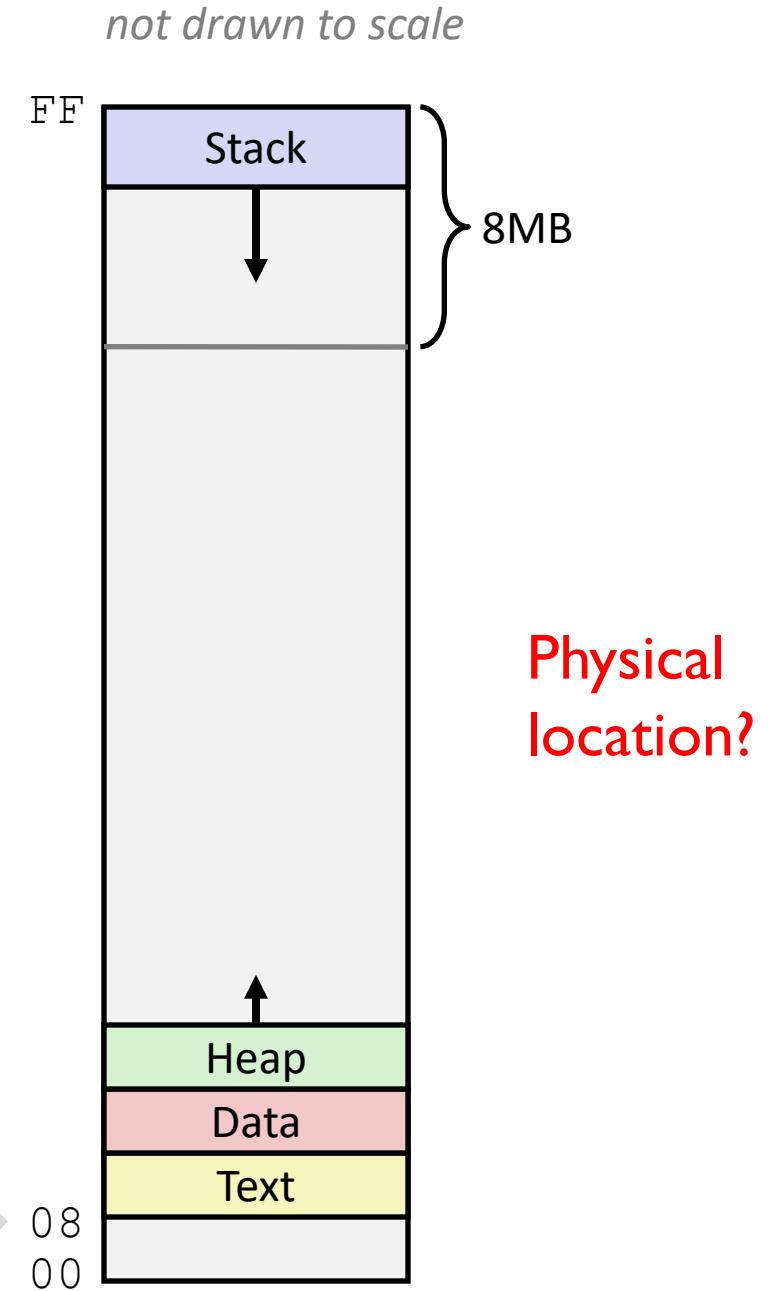
Locality

CPSC 275
Introduction to Computer Systems

IA-32 Linux Memory Layout

- Stack
 - Runtime stack (8MB limit)
 - e. g., local variables
- Heap
 - Dynamically allocated storage
 - When call malloc, calloc, new
- Data
 - Statically allocated data
 - e.g., globals & strings declared in code
- Text
 - Executable machine instructions
 - Read-only (*re-entrant*)

Upper 2 hex digits
= 8 bits of address



Random-Access Memory (RAM)

- Key features
 - RAM is traditionally packaged as a chip.
 - Basic storage unit is normally a cell (one bit per cell).
 - Multiple RAM chips form a memory.
- RAM comes in two varieties:
 - DRAM (dynamic RAM)
 - SRAM (static RAM)



SRAM vs DRAM

	Trans. per bit	Access time	Needs refresh?	Needs EDC?	Cost	Applications
SRAM	4 or 6	IX	No	Maybe	100x	Cache memories
DRAM	I	10X	Yes	Yes	IX	Main memories, frame buffers

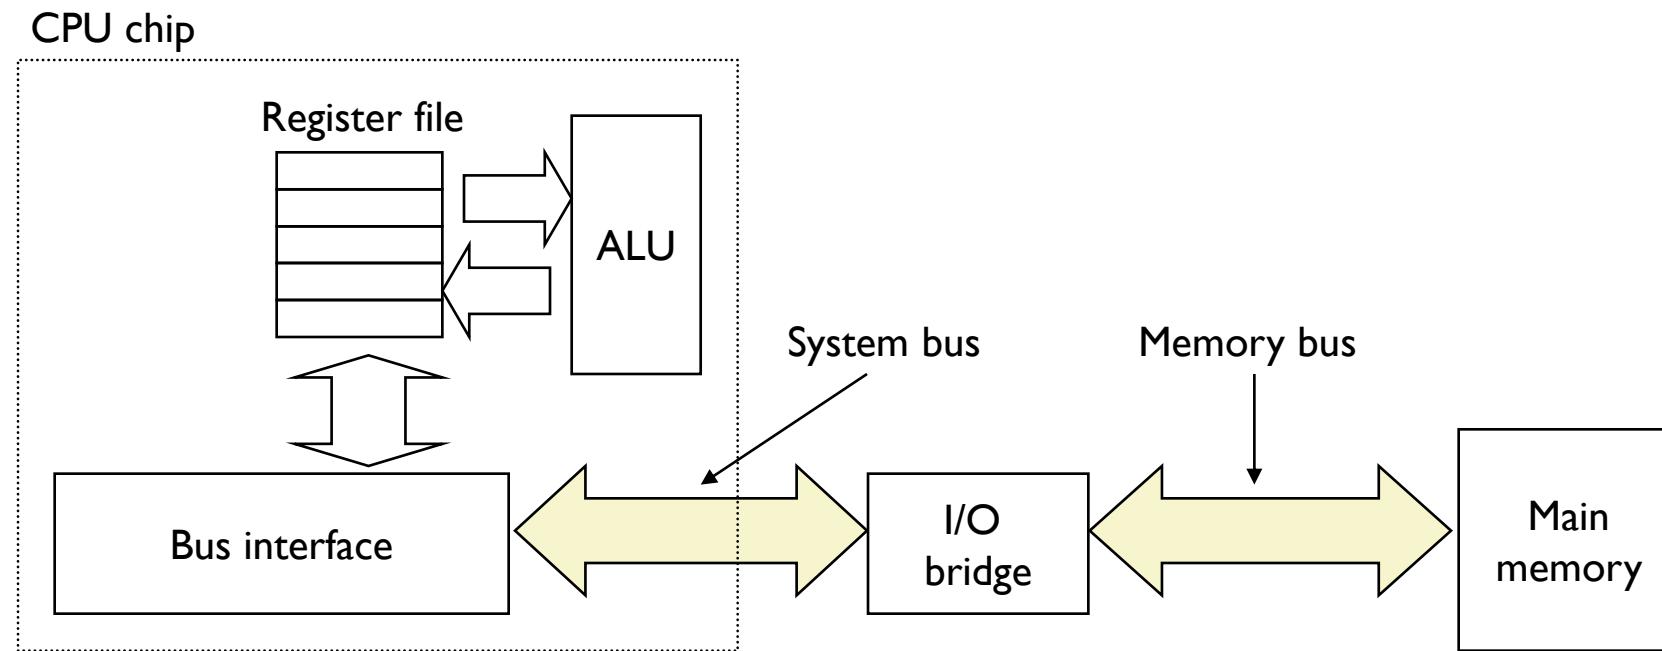
DRAM and SRAM are *volatile* memories
(lose information if powered off)

Nonvolatile Memories

- *Nonvolatile* memories retain value even if powered off
 - Read-only memory (**ROM**): programmed during production
 - Programmable ROM (**PROM**): can be programmed once
 - Eraseable PROM (**EPROM**): can be bulk erased (UV, X-Ray)
 - Electrically eraseable PROM (**EEPROM**): electronic erase capability
 - Flash memory: EEPROMs. with partial (block-level) erase capability
- Uses for nonvolatile memories
 - *Firmware* programs stored in a ROM (BIOS, controllers for disks, network cards, graphics accelerators, security subsystems,...)
 - Solid state disks (SSD)
 - Disk caches

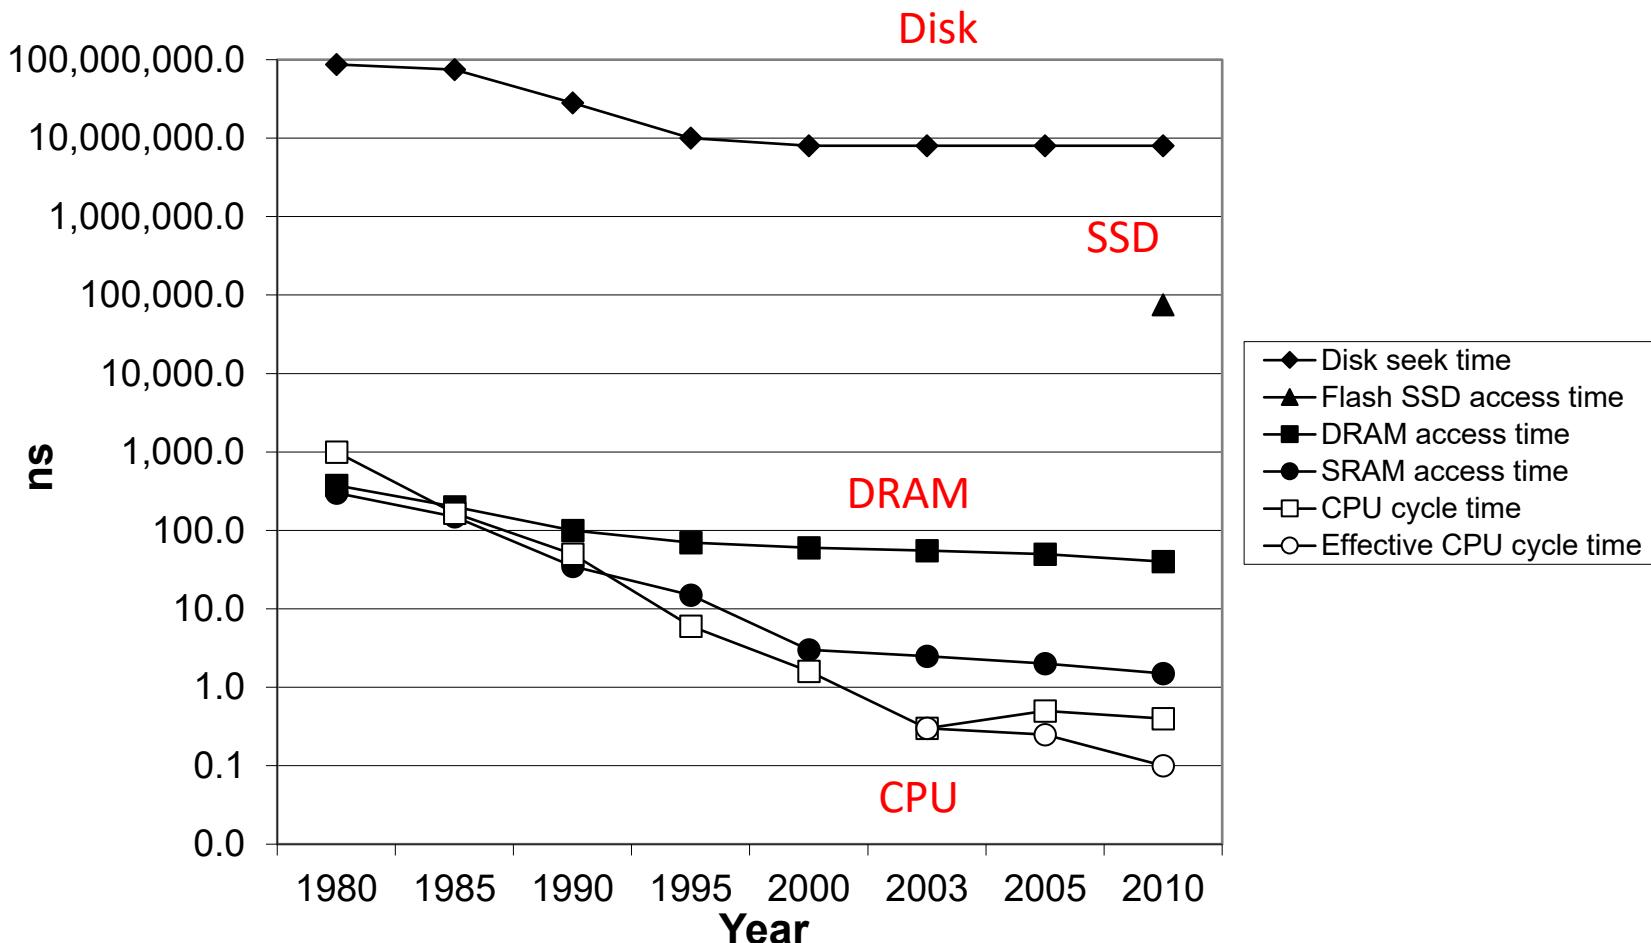
Connecting CPU and Memory

- A **bus** is a collection of parallel wires that carry address, data, and control signals.
- Buses are typically shared by multiple devices.



The CPU-Memory Gap

The gap widens between DRAM, disk, and CPU speeds.

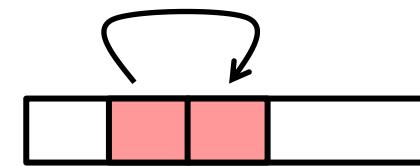
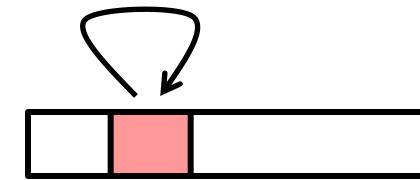


Locality to the rescue!

The key to bridging this CPU-Memory gap is a fundamental property of computer programs known as **locality**.

What is Locality?

- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently
- **Temporal locality:**
 - Recently referenced items are likely to be referenced again in the near future
- **Spatial locality:**
 - Items with nearby addresses tend to be referenced close together in time



Locality Example

```
sum = 0;  
for (i = 0; i < n; i++)  
    sum += a[i];  
return sum;
```

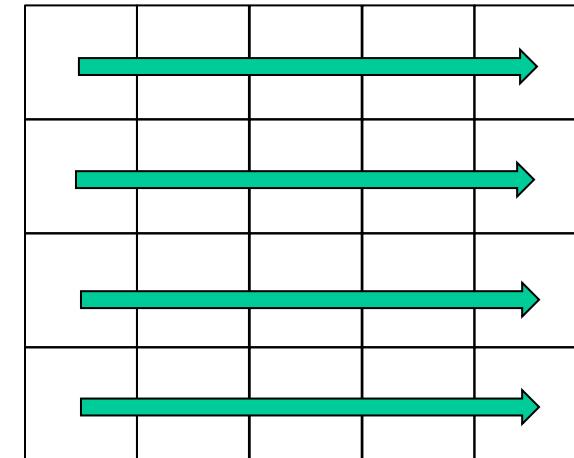
- data references
 - Reference array elements in succession (stride-1 reference pattern). spatial locality
 - Reference variable **sum** each iteration. temporal locality
- instruction references
 - Reference instructions in sequence. spatial locality
 - Cycle through loop repeatedly. temporal locality

Qualitative Estimates of Locality

- Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer.
- **Q:** Does this function exhibit good locality with respect to array *a*?

```
int sum_array_rows(int a[M] [N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i] [j];
    return sum;
}
```



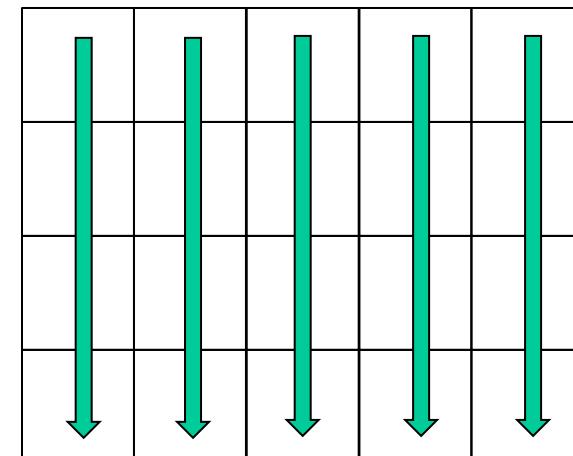
Q: Stride? 1

Locality Example

- **Q:** Does this function exhibit good locality with respect to array a?

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```



Q: Stride? N

- **Q:** Which one is better?

Locality Example

- **Q:** Can you permute the loops so that the function scans the 3-d array `a` with a stride-1 reference pattern and thus has good spatial locality? (Homework)

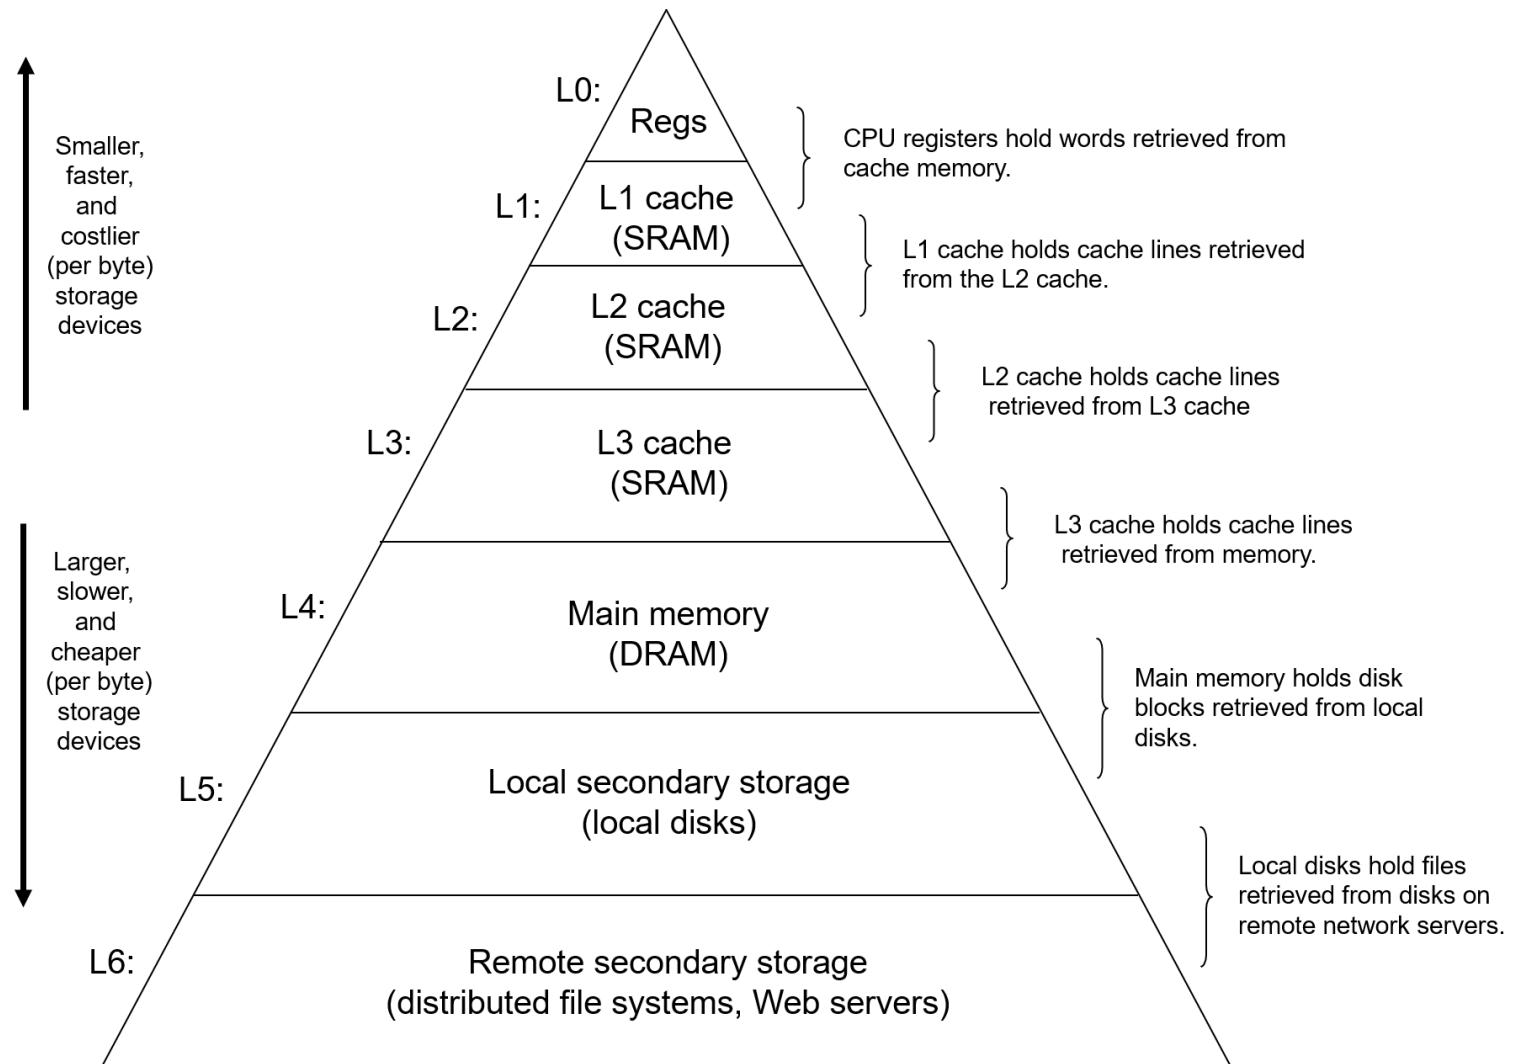
```
int sum_array_3d(int a[N] [N] [N] )
{
    int i, j, k, sum = 0;

    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            for (k = 0; k < N; k++)
                sum += a[k] [i] [j];
    return sum;
}
```

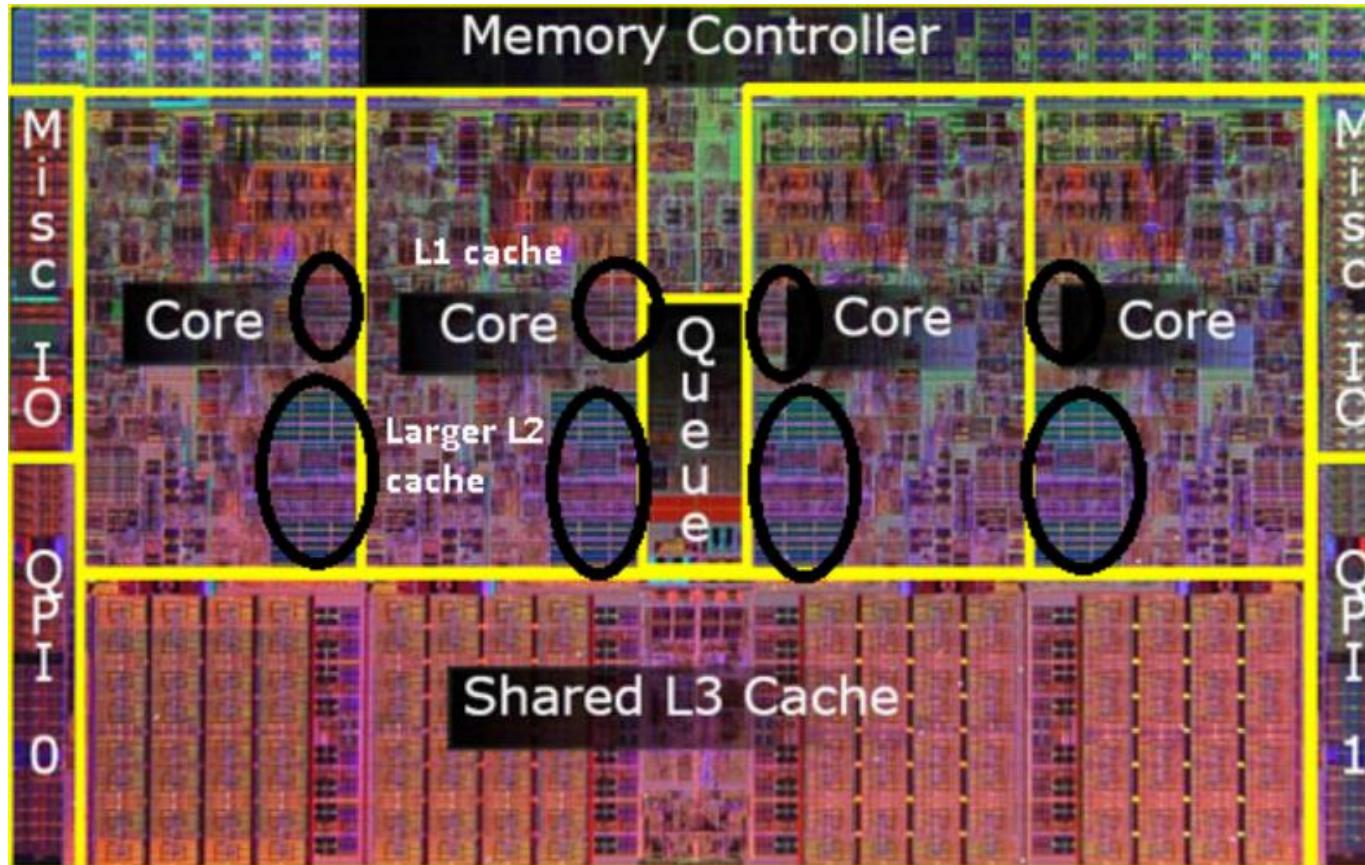
Memory Hierarchy

- Some fundamental and enduring properties of hardware and software:
 - Fast storage technologies cost more per byte, have less capacity, and require more power (heat!).
 - The gap between CPU and main memory speed is widening.
 - **Well-written** programs tend to exhibit **good locality**.
- This led to an approach of organizing memory and storage systems known as a **memory hierarchy**.

Memory Hierarchy



Where is cache?



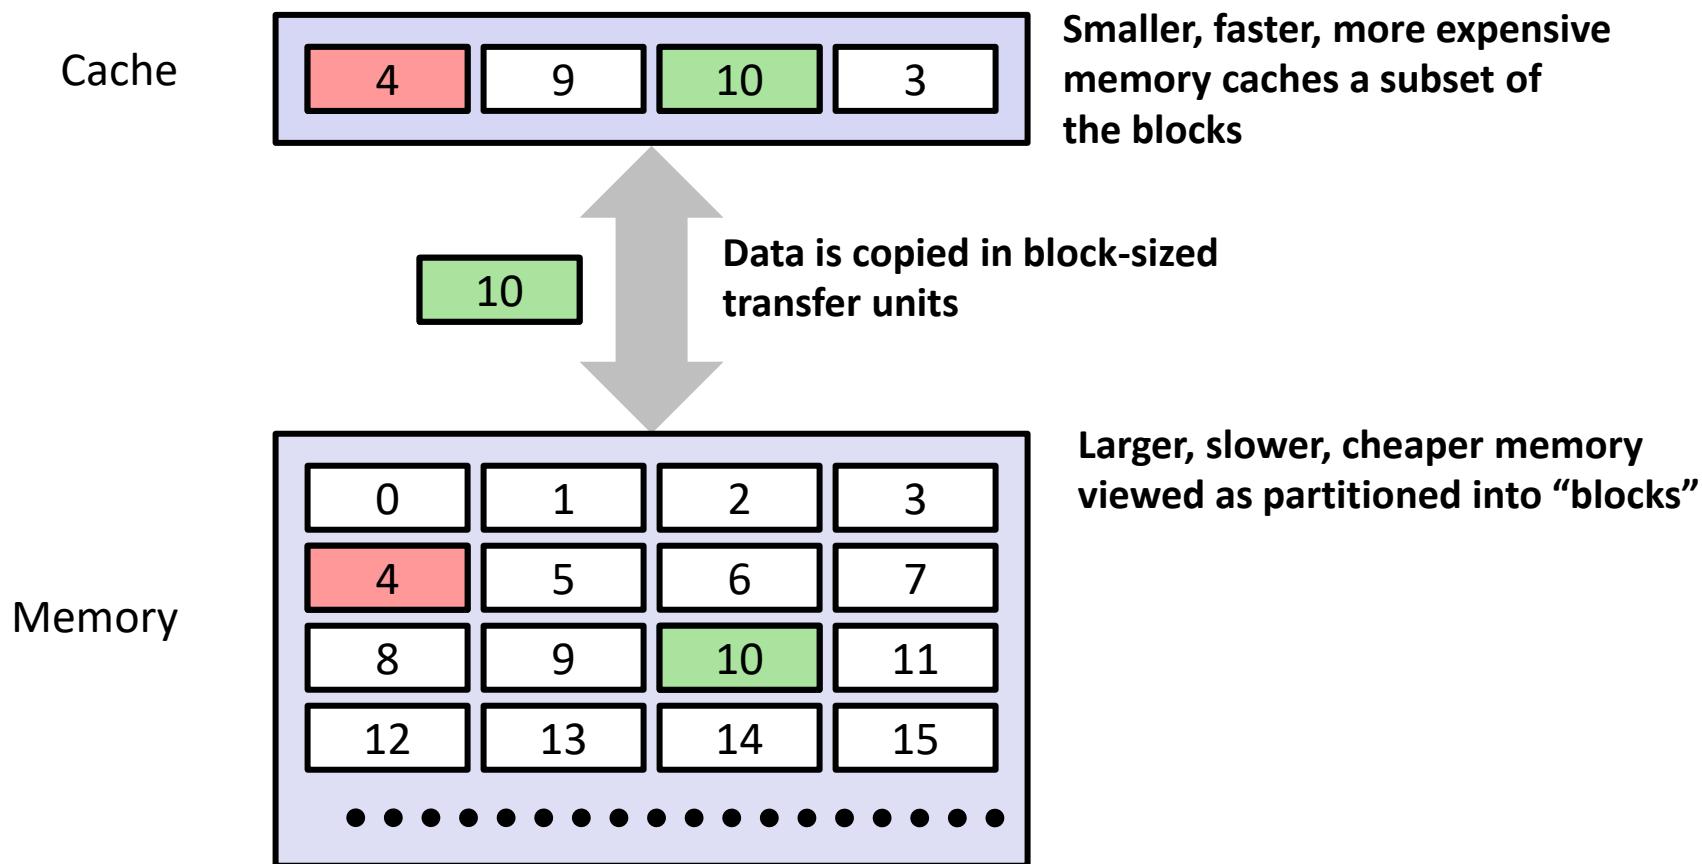
A typical quad-core processor (Intel i7)

Cache

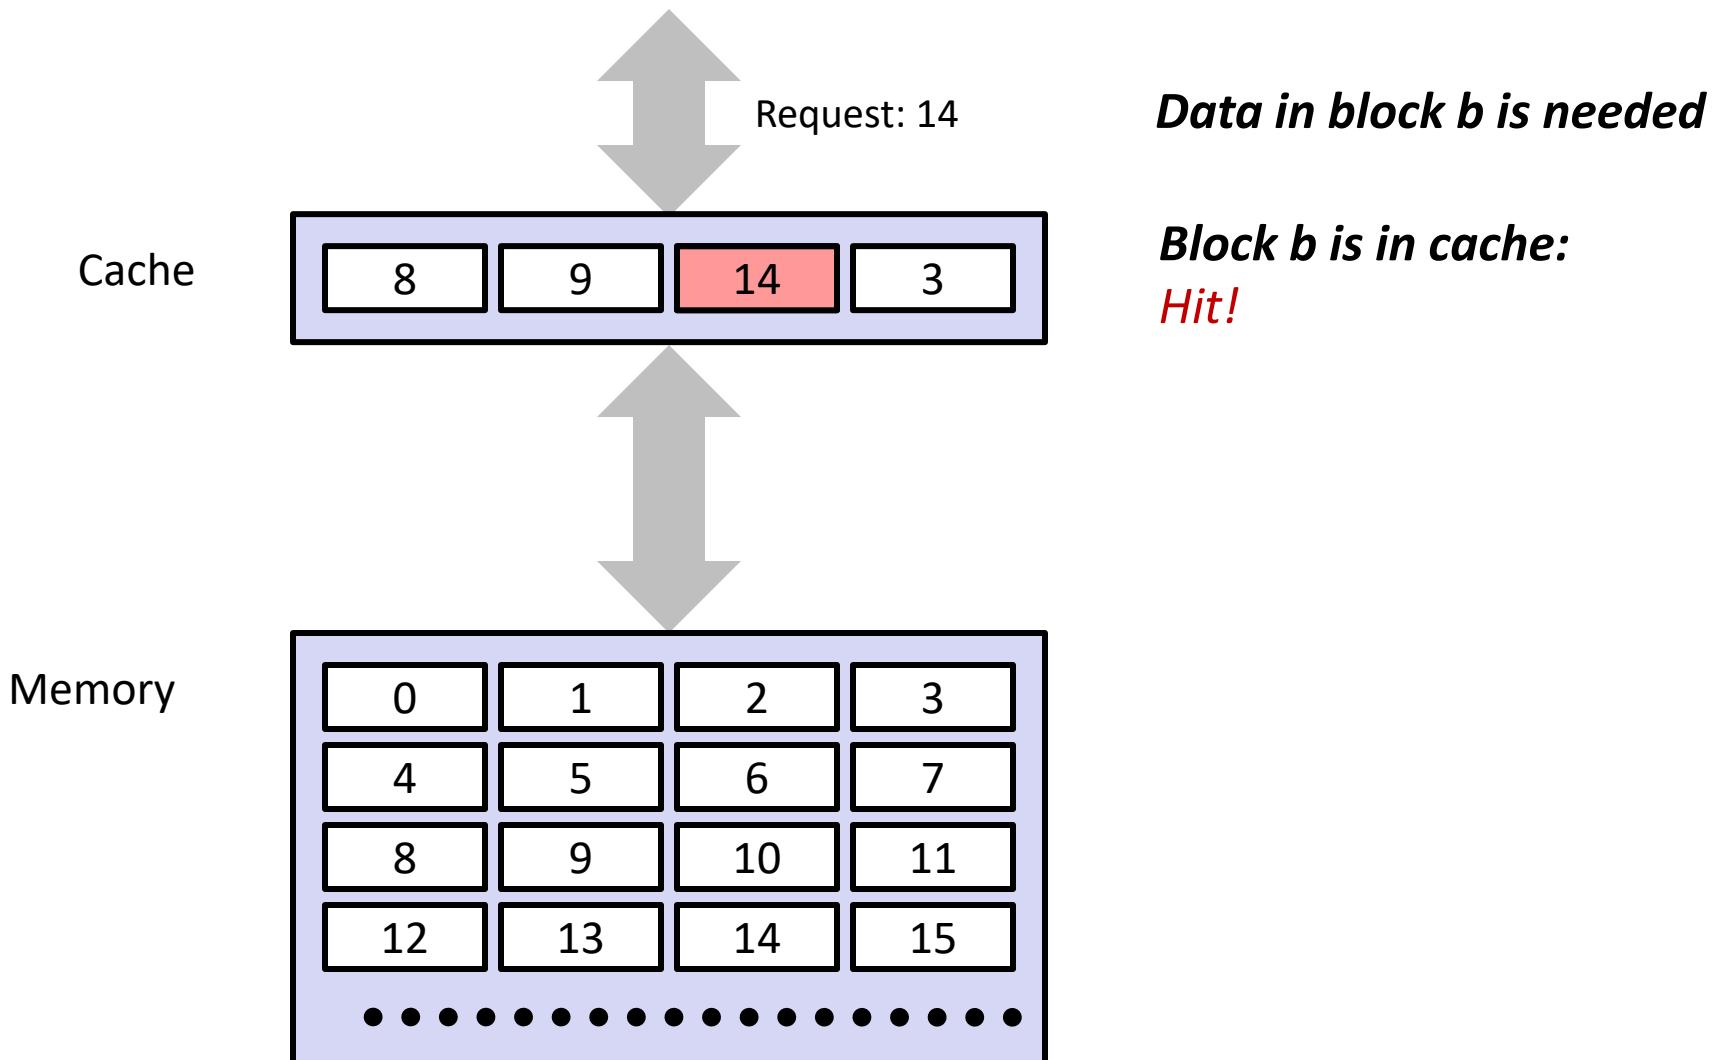
- **Cache:** A smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device.
- Fundamental idea of a memory hierarchy:
 - For each k , the faster, smaller device at level k serves as a cache for the larger, slower device at level $k+1$.
- Why do memory hierarchies work?
 - Because of locality, programs tend to access the data at level k more often than they access the data at level $k+1$.
 - Thus, the storage at level $k+1$ can be slower, and thus larger and cheaper per bit.



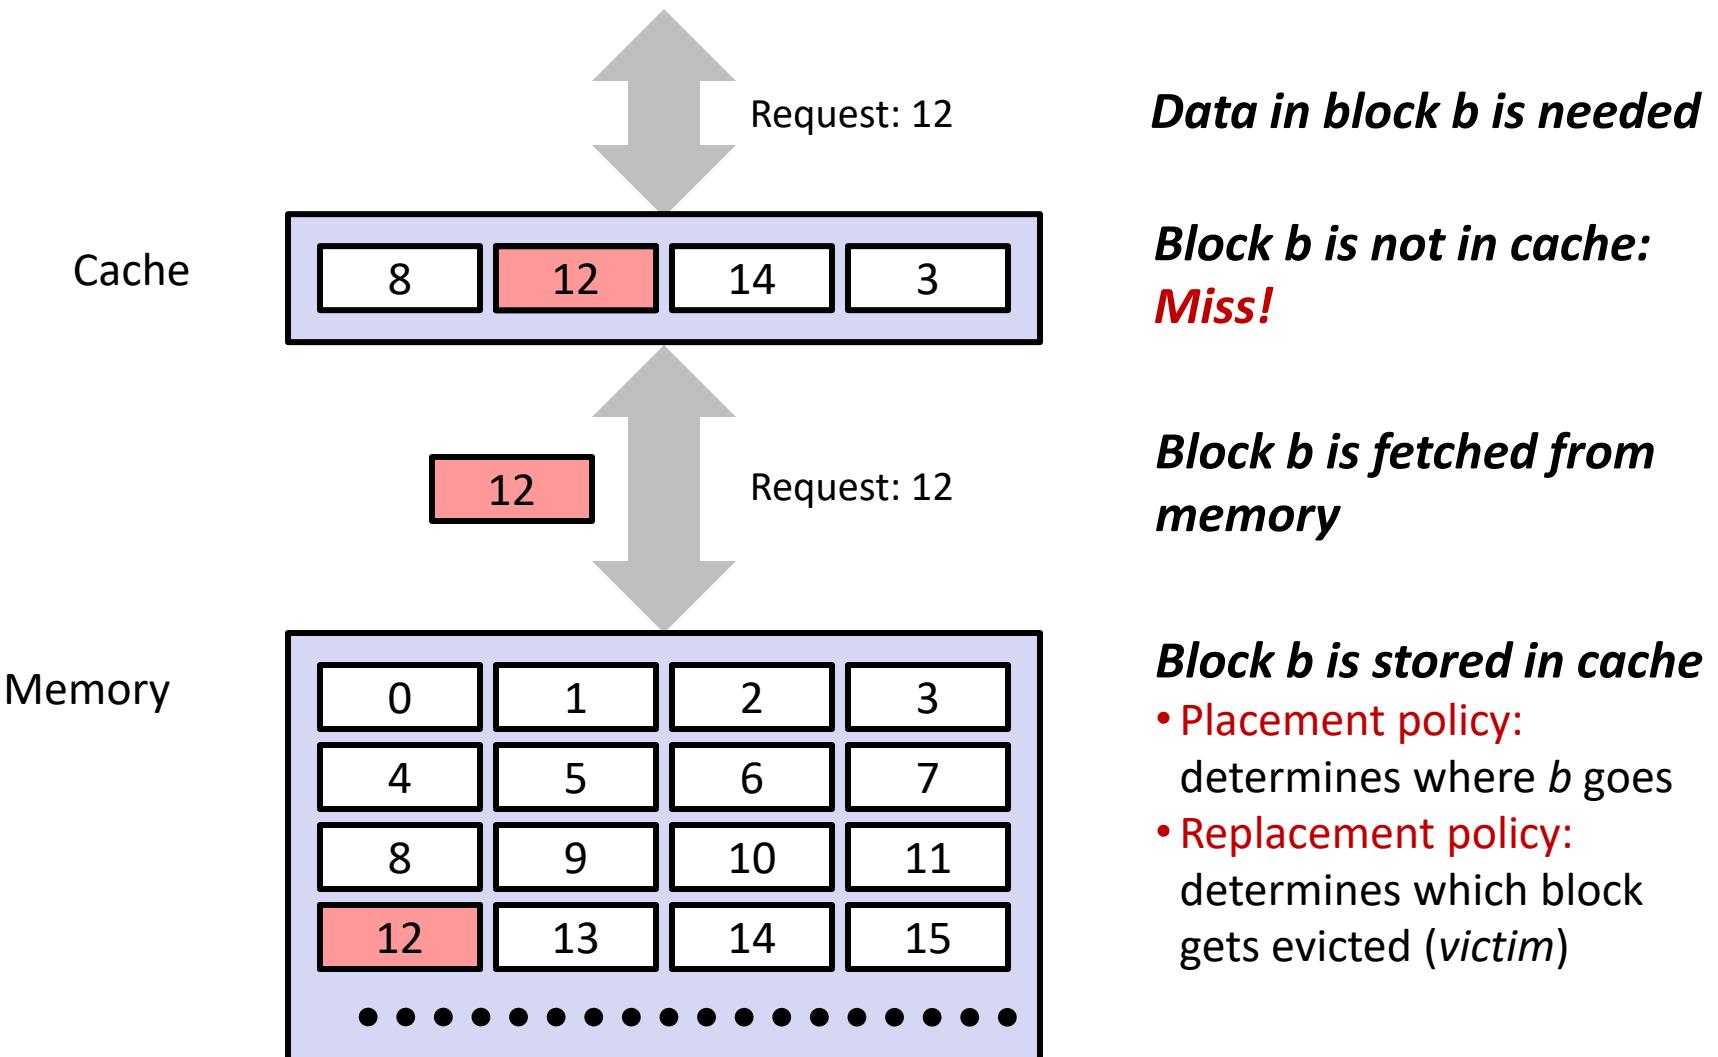
General Cache Concepts



General Cache Concepts



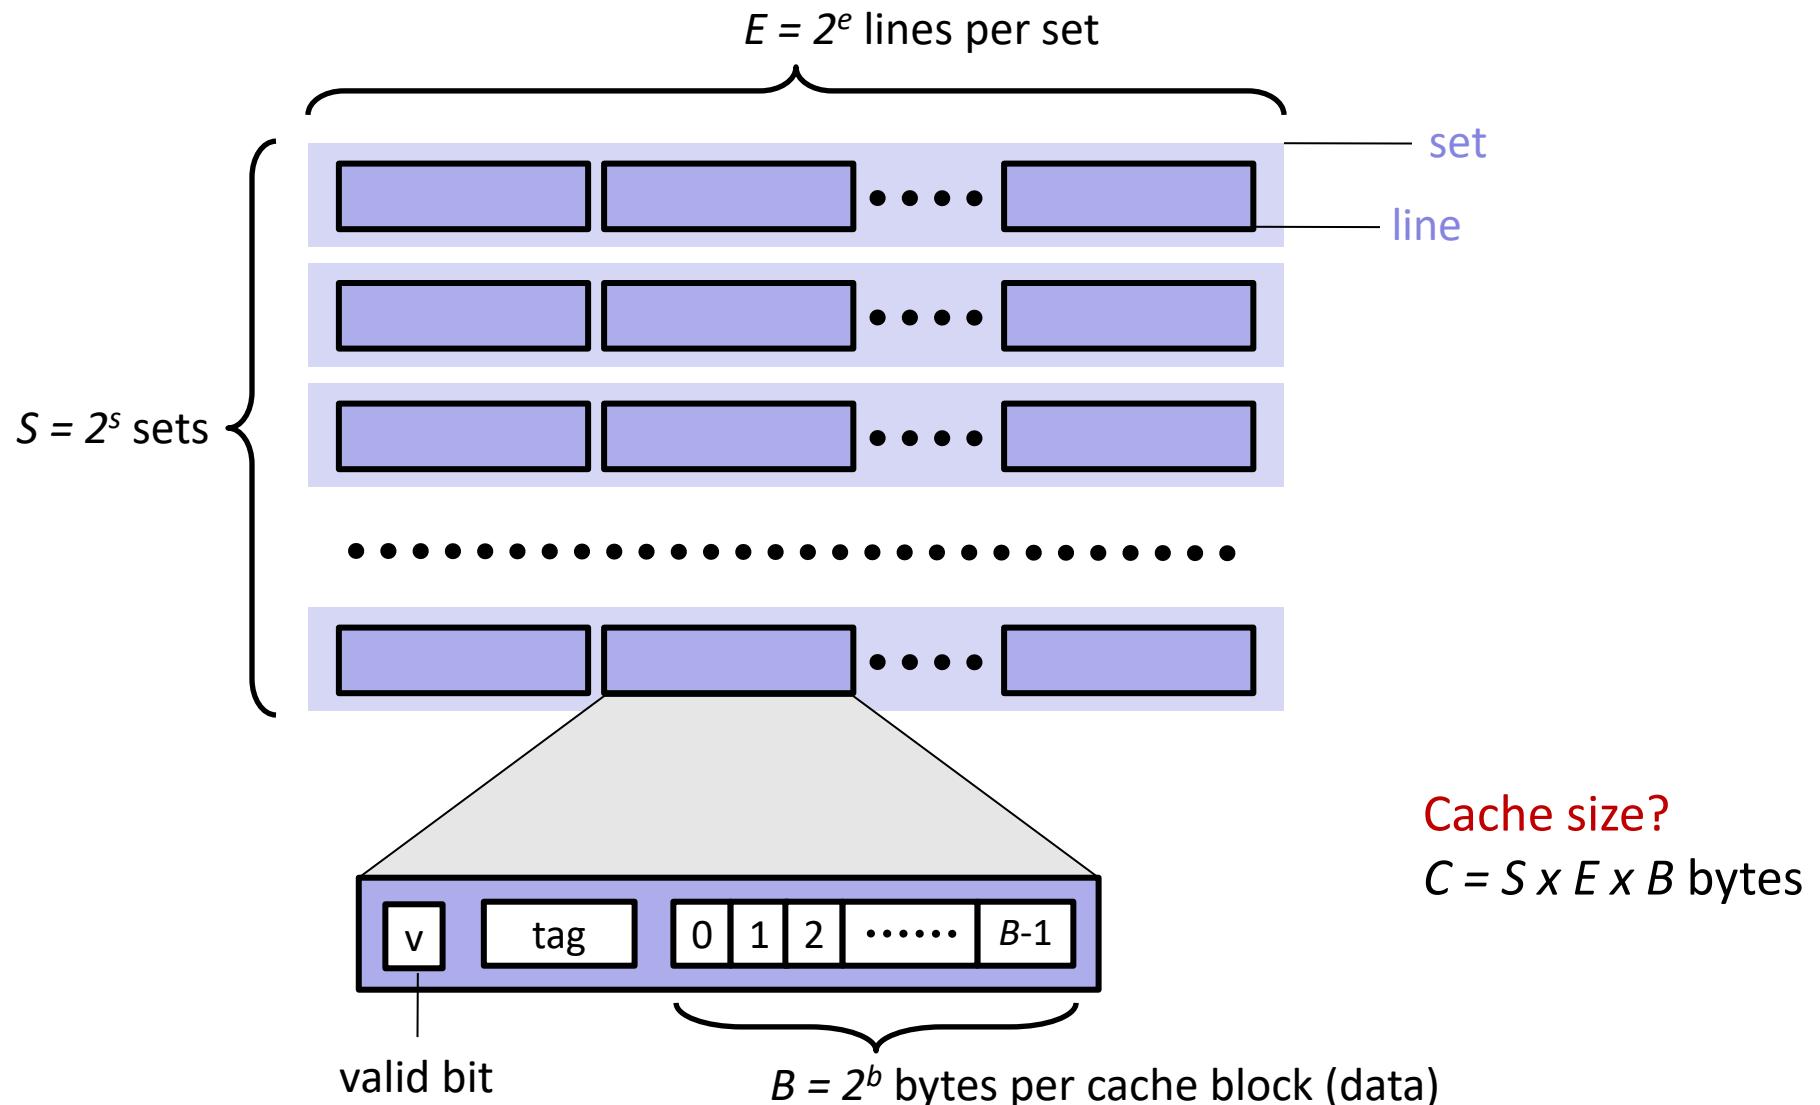
General Cache Concepts: Misses



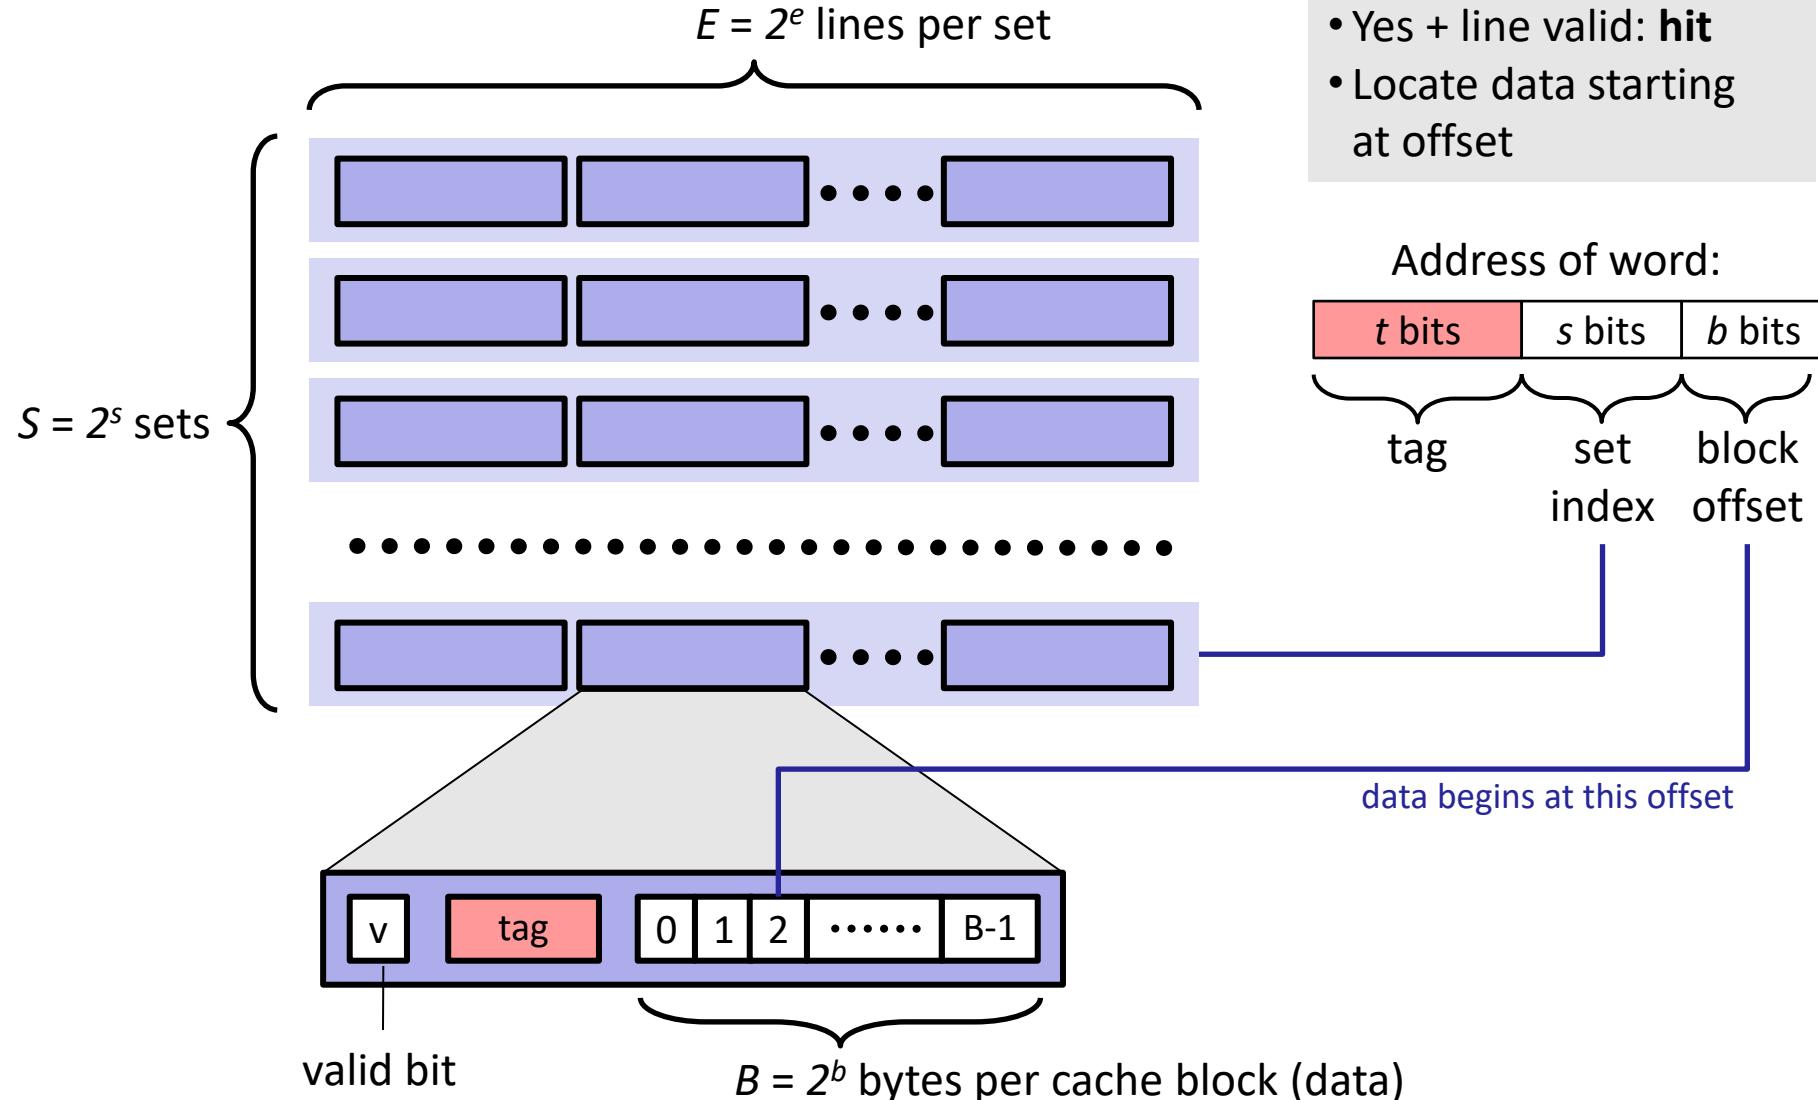
Types of Cache Misses

- **Cold (compulsory) miss**
 - Cold misses occur because the cache is empty.
- **Conflict miss**
 - Most caches limit blocks at level $k+1$ to a small subset of the block positions at level k .
 - Conflict misses occur when the level k cache is large enough, but multiple data objects all map to the same level k block.
 - e.g., Referencing blocks 0, 8, 0, 8, 0, 8, ... would miss every time.
- **Capacity miss**
 - Occurs when the set of active cache blocks (known as the *working set*) is larger than the cache.

General Cache Organization (S, E, B)



Cache Read

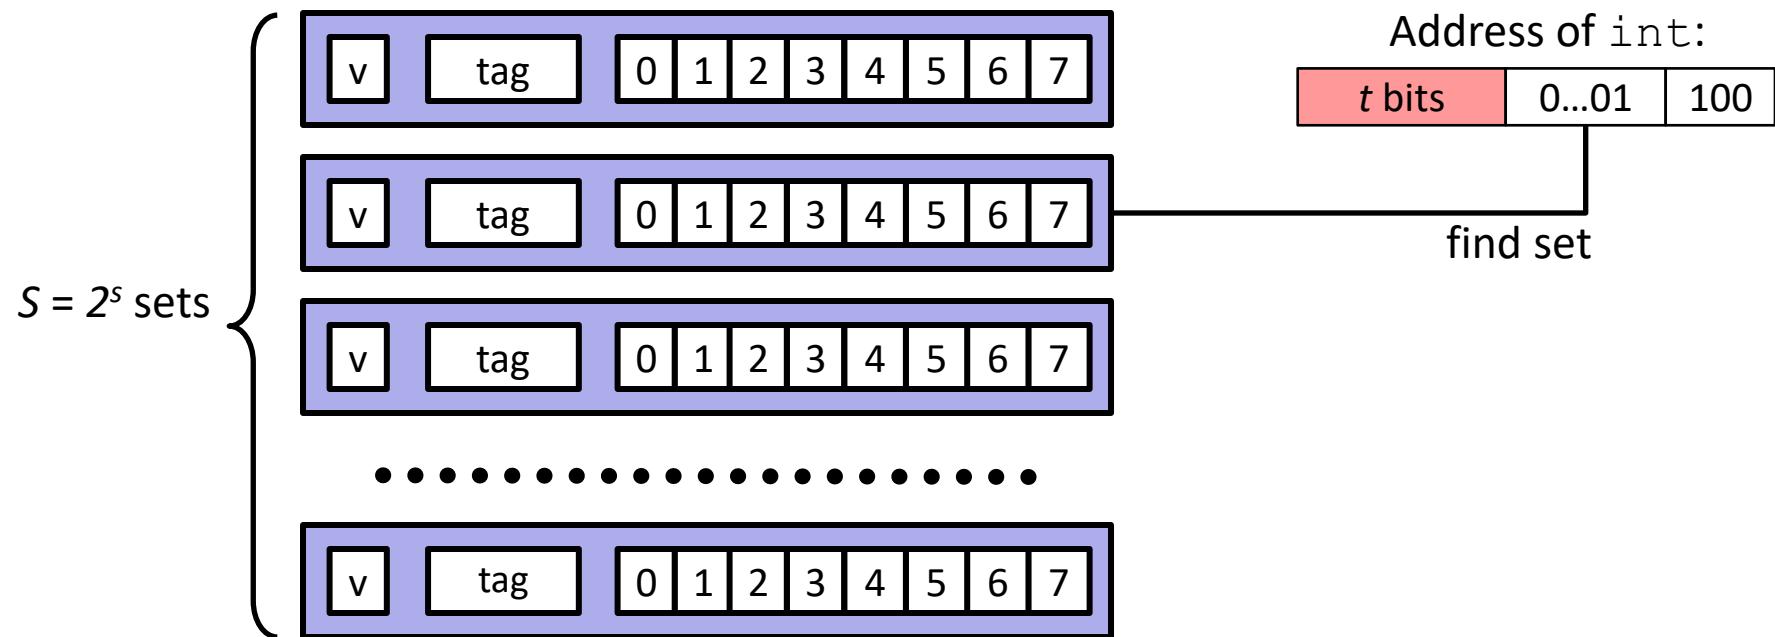




Example: Direct Mapped Cache ($E = 1$)

Direct mapped: One line per set

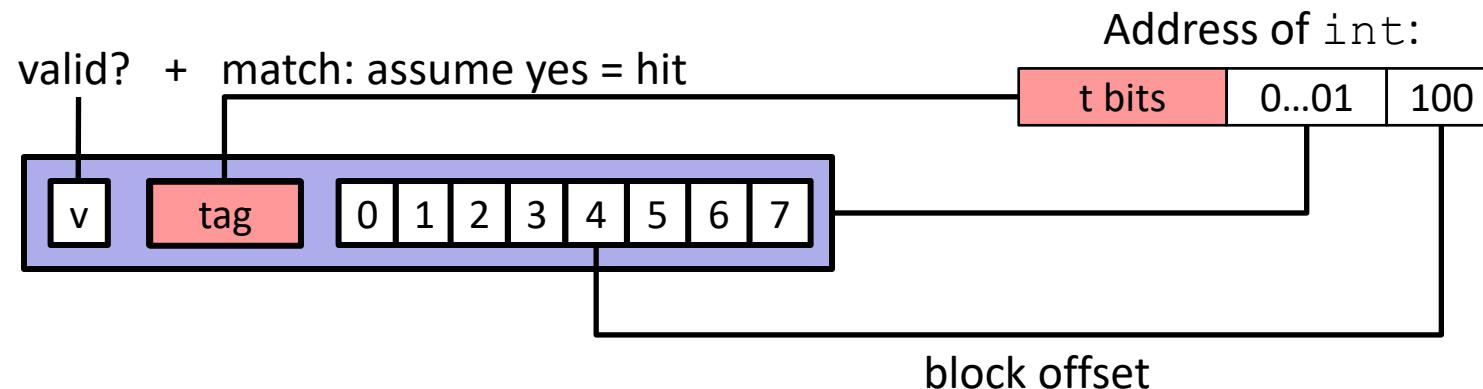
Assume: cache block size 8 bytes



Example: Direct Mapped Cache ($E = 1$)

Direct mapped: One line per set

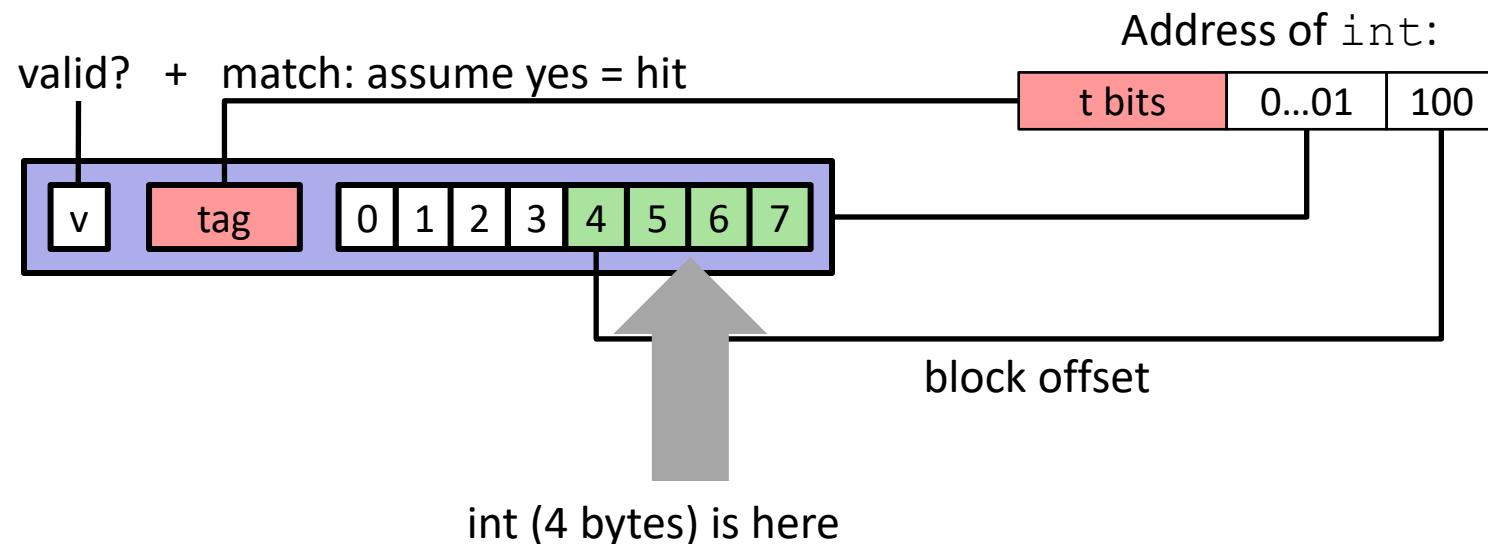
Assume: cache block size 8 bytes



Example: Direct Mapped Cache ($E = 1$)

Direct mapped: One line per set

Assume: cache block size 8 bytes



If no match, an old line is evicted and replaced.

Direct-Mapped Cache Simulation

4-bit address

x	xx	x
---	----	---

t=1 s=2 b=1

M=16 byte addresses (total size of memory)

B=2 bytes/block

S=4 sets

E=1 line/set

Direct-Mapped Cache Simulation

4-bit address

x	xx	x
---	----	---

t=1 s=2 b=1

M=16 byte addresses (total size of memory)

B=2 bytes/block

S=4 sets

E=1 line/set

Address trace:

0	[0000 ₂]	miss
1	[0001 ₂]	hit
7	[0111 ₂]	miss
8	[1000 ₂]	miss
0	[0000 ₂]	miss

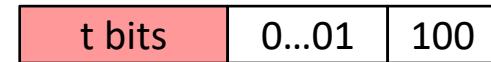
	v	Tag	Block
Set 0	1	0	M[0-1]
Set 1			
Set 2			
Set 3	1	0	M[6-7]

2-way Set Associative Cache

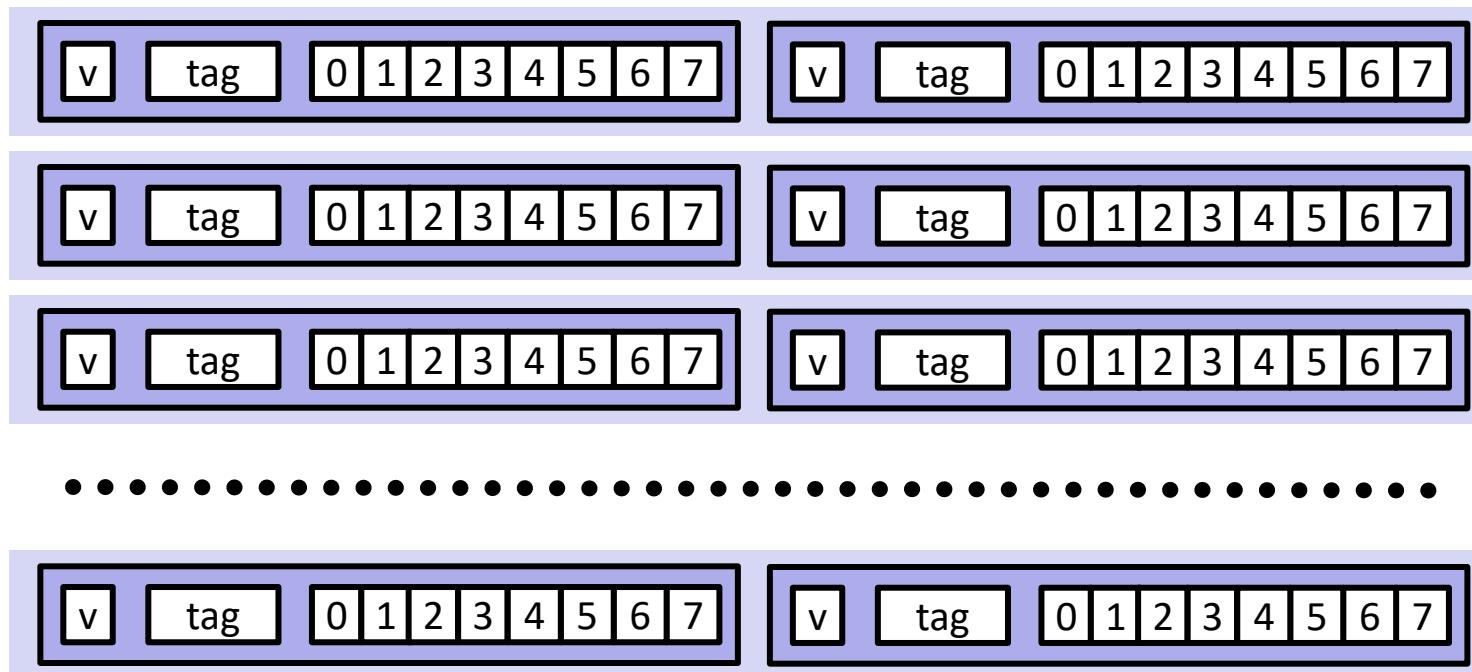
$E = 2$: Two lines per set

Assume: cache block size 8 bytes

Address of short int:



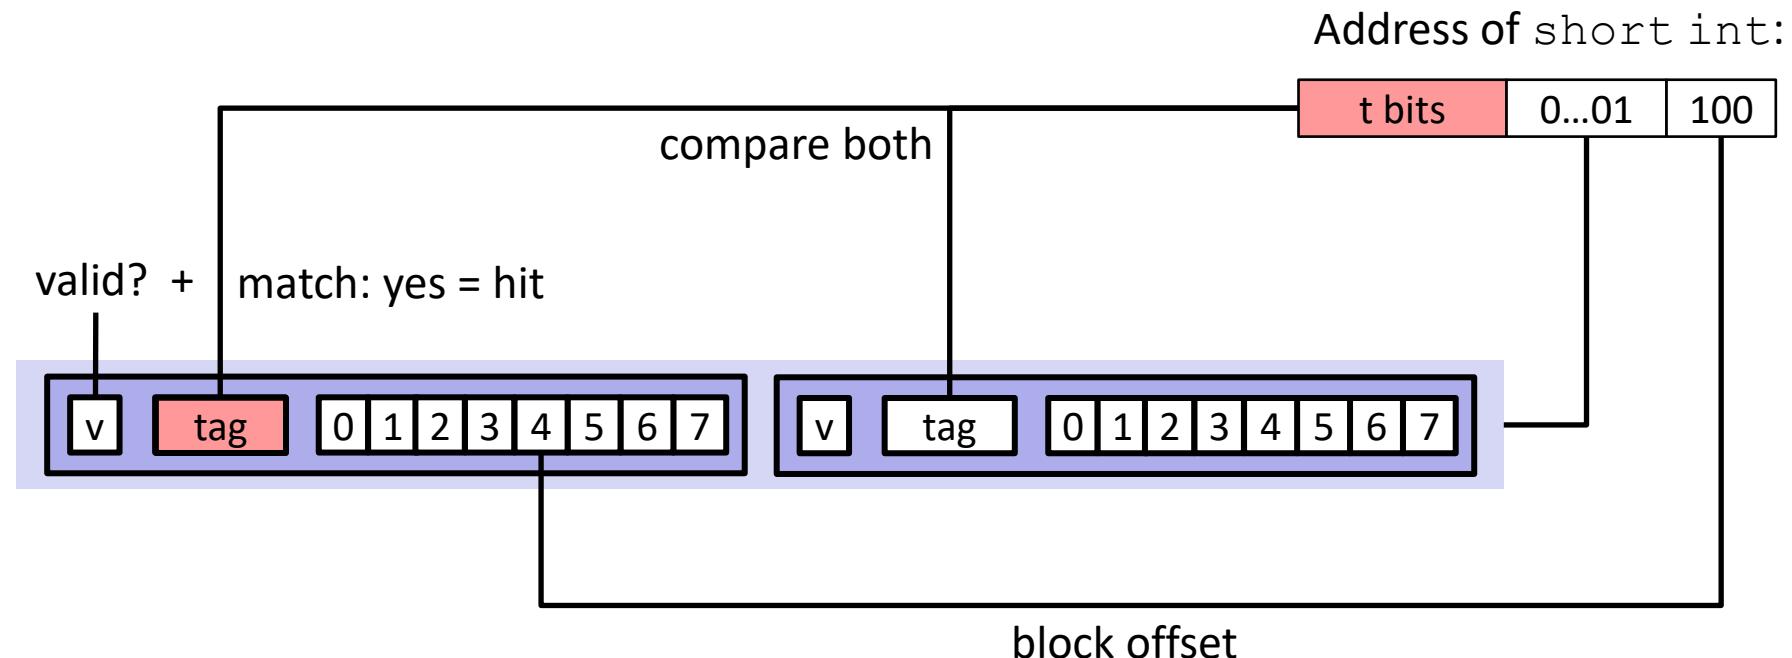
find set



2-way Set Associative Cache

$E = 2$: Two lines per set

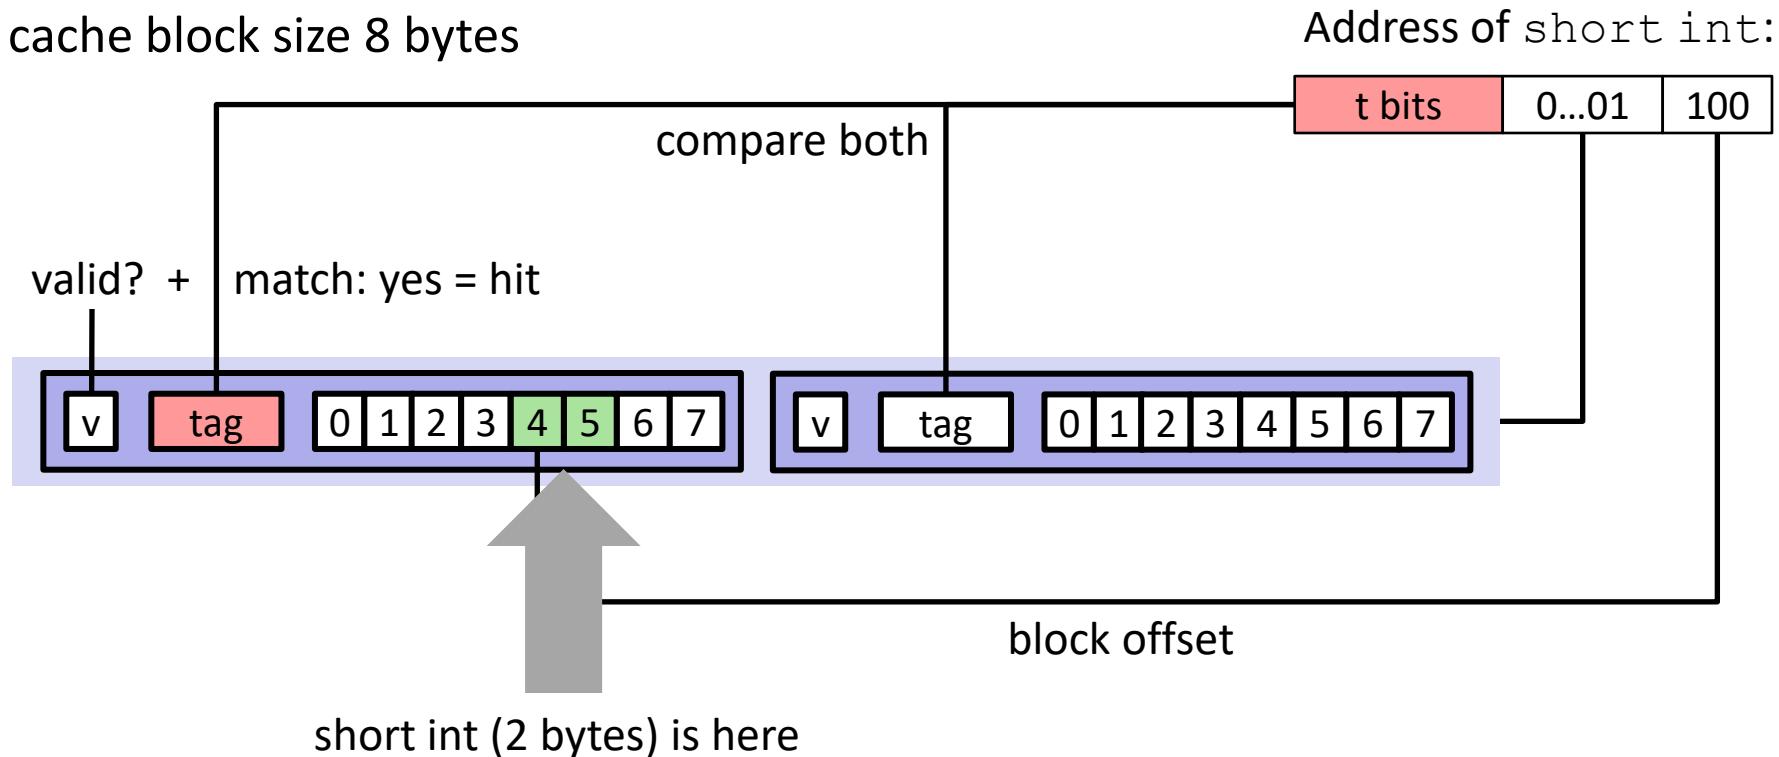
Assume: cache block size 8 bytes



2-way Set Associative Cache

E = 2: Two lines per set

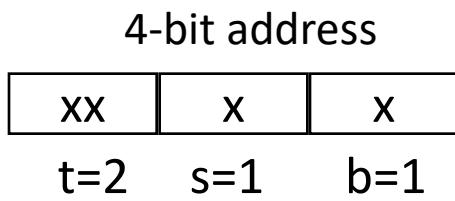
Assume: cache block size 8 bytes



No match:

- One line in set is selected for eviction and replacement
- Replacement policies: random, least recently used (LRU), ...

2-Way Set Associative Cache Simulation



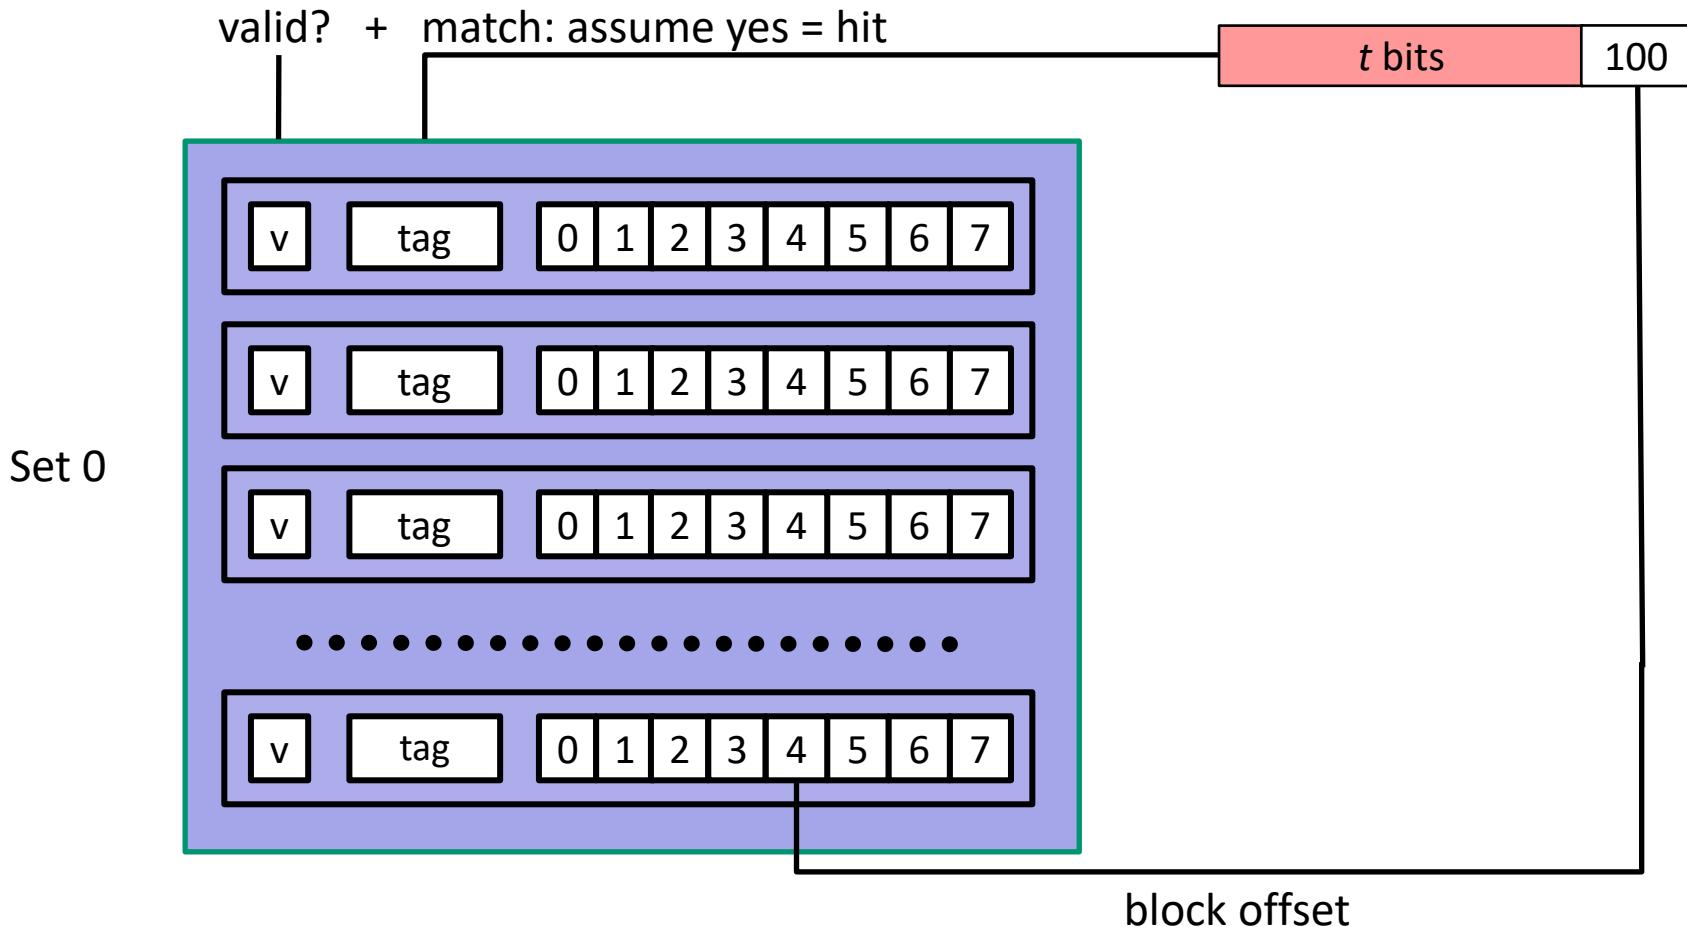
M=16 byte addresses,
B=2 bytes/block,
S=2 sets,
E=2 lines/set

Address trace (reads, one byte per read):

0	[0000 ₂]	miss
1	[0001 ₂]	hit
7	[0111 ₂]	miss
8	[1000 ₂]	miss
0	[0000 ₂]	hit

	v	Tag	Block
Set 0	1	00	M[0-1]
	1	10	M[8-9]
Set 1	1	01	M[6-7]
	0		

Fully Associative Cache ($S = 1$)





Assignment 7: Cache Simulator

- A C program that simulates the behavior of a cache memory
 - Parameters: s , E , b
 - Given a memory access *trace file* as input, simulates the hit/miss behavior of a cache memory on this trace and prints the total number of hits, misses, and evictions.

Trace Files

- Trace files generated by valgrind, a Linux utility program for memory debugging, memory leak detection, and profiling
 - Example trace file

I 0400d7d4, 8	// instruction load
M 0421c7f0, 4	// modify (a load followed by a store)
L 04f6b868, 8	// load
S 7ff0005c8, 8	// store

↑ ↑ ↑
space address data size

Trace Files, cont'd

- For this assignment, we are interested only in data cache performance, so your simulator should ignore all instruction cache accesses (lines starting with “I”).
- To help you simplify the code, all lines with an instruction load have been already removed from trace files.

LRU Replacement Policy

- Necessary conditions for replacement
 - a cache miss occurred
 - the current set is full
- Need to determine which line from the current set will be replaced by a new line (*eviction*).
- Pick the one *least recently used* (referenced).
 - Why does it work?
 - How to implement?

Running Your Simulator

```
$ ./mycache -s 2 -E 1 -b 4 -t traces/tiny.trace
```

To verify the result:

```
$ ./refcache -s 2 -E 1 -b 4 -t traces/tiny.trace
```

Programming Notes

- For arbitrary s , E , and b , you will need to allocate storage for your simulator's data structures using `malloc`.
- 64-bit address field: `unsigned long int`
- Ignore the request data sizes in the traces, that is, there is no need to define blocks in lines.
- Use a timestamp to keep track of the last time each line has been referenced.
- Start with small trace files.



Command-Line Arguments

- When we run a program, we'll often need to supply it with information.
- Example:

```
$ ./repeat 10 computer
```

Command-Line Arguments

- To obtain access to *command-line arguments*, main must have two parameters:

```
int main(int argc, char *argv[] )  
{  
    ...  
}
```

Command-Line Arguments

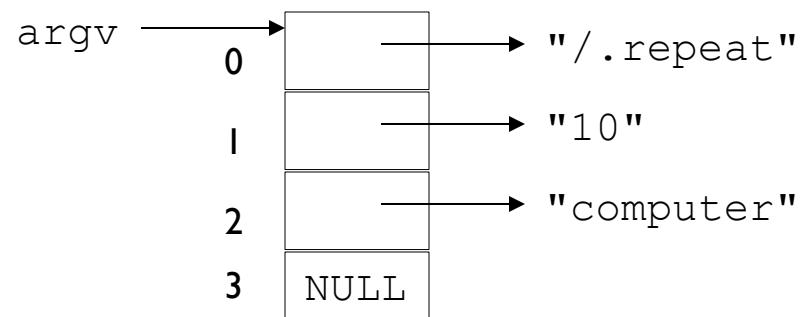
- `argc` (“argument count”) is the number of command-line arguments.
- `argv` (“argument vector”) is an array of pointers to the command-line arguments (stored as strings).
- `argv[0]` points to the name of the program,
- `argv[1]` through `argv[argc-1]` point to the remaining command-line arguments.

Command-Line Arguments

- If the user enters the command line

```
$ ./repeat 10 computer
```

then argc will be 3 (why?),
and argv will have the following appearance:



Command-Line Arguments

- Since `argv` is an array of pointers, accessing command-line arguments is easy, e.g,

```
int i;  
for (i = 1; i < argc; i++)  
    printf("%s\n", argv[i]);
```

Output:

```
./repeat  
10  
computer
```

- For numeric arguments, conversion might be necessary:

```
int count = atoi(argv[1]);
```

Command-Line Options

- Command-line *options* modify the program's behavior.
- Command-line arguments and options

```
$ ls -l myfile
```

option argument

- Some command-line options take values:

```
$ tail -n 20 myfile
```

- The **getopt** function is useful in parsing command-line options.

The getopt Function

```
#include <unistd.h>

int getopt(int argc, char *argv[], char *optstring);
extern char *optarg;
```

- **argc** is the number of arguments.
- **argv** is an array of arguments.
- **optstring** is a string containing the option characters.
 - If such a character is followed by a colon, the option requires a value.
- **optarg** is an *external* variable string containing the option values.
- If there are no more option characters, the function returns -1.

The getopt Function

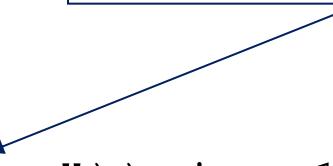
- Normally, getopt is called in a loop.
- When getopt returns -1, indicating no more options are present, the loop terminates.
- A switch statement is used to dispatch on the return value from getopt.

Example using getopt

```
#include <stdio.h>
#include <unistd.h>

int main (int argc, char *argv[]) {
    int r, cval;
    while ((r = getopt(argc, argv, "ab:c:")) != -1)
        switch (r) {
            case 'a':
                printf("Do option a!\n"); break;
            case 'b':
                printf("Do option b with %s!\n", optarg); break;
            case 'c':
                cval = atoi(optarg);
                printf("Do option c with %d!\n", cval);
                break;
            default:
                printf("Error: Unknown option!\n");
        }
}
```

three options: a, b, and c
b and c require values
c's value is a number



Opening a file

```
#include <stdlib.h>
FILE *fopen(char *filename, char *mode);
```

- *filename* is the name of the file to be opened.
- *mode* is a “mode string” that specifies what operations we intend to perform on the file.
- returns a *file pointer* (or **NULL** on error):

```
FILE fp = fopen("data.in", "r");
... // do something with the file
fclose(fp);
```

Modes

Mode strings for text files:

String	Meaning
"r"	Open for reading
"w"	Open for writing (file need not exist)
"a"	Open for appending (file need not exist)
"r+"	Open for reading and writing, starting at beginning
"w+"	Open for reading and writing (truncate if file exists)
"a+"	Open for reading and writing (append if file exists)

Reading from a file

```
char *fgets(char *buf, int n, FILE *stream)
```

- *buf* is the pointer to an array of chars where the string read is stored.
- *n* is the maximum number of characters to be read, including the null character.
- *stream*, is the pointer to a FILE object
- Example:

```
FILE fp = fopen("data.in", "r");  
char buf[10];  
fgets(buf, 10, fp);
```

Reading from a string

```
int sscanf(char *str, char *format, ...)
```

- *str* is the pointer to an array of chars where the string is stored.
- *format* is a format specifier
- Example:

```
char s[] = "November 20, 2023";
char mon[20];
int day, yr;
sscanf(s, "%s %d, %d", mon, &day, &yr);
```

→ mon = “November”, day = 20, yr = 2023

