

# Announcements


- Exam 3
  - Friday, November 21
  - Covers up to Cache: Lectures 19-28
  - Format: multiple-choice (30%) / short-answer (70%)
- Assignment 8
  - Posted Thursday, November 20; due December 2
- Quiz 11
  - Monday, December 1
  - Covers Lectures 29-31
- Graded Lab 3
  - December 3-4
  - Covers Assignments 4-7

Lecture 30

# Systems Programming

CPSC 275  
Introduction to Computer Systems

# A typical computer system



# The Linux System

- *Kernel* – heart of the OS (always running)
  - Process scheduling
  - Memory management
  - I/O control
  - many more ...
- *Shell* – Interpreter between the user and the computer
  - e.g., bash
- Tools and applications
  - Accessible from shell
  - Can be run independently of shell

# Shell

- Provides command line as an interface between the user and the system
- Starts automatically when you log in
  - e.g., bash for Ubuntu Linux
- Uses a command *language*
  - Allows programming (shell scripting) within the shell environment
    - Uses variables, loops, conditionals, etc.
  - Accepts commands and often makes *system calls* to carry them out

# What is Systems Programming?

- Application Programming
  - Solve a user-facing problem (e.g., web browser, text editor).
  - High-level (e.g., Python, Java, JavaScript). Hides machine details.
  - User interface and features.
- Systems Programming
  - Provide services for other software to run (e.g., OS, drivers, compilers, shells).
  - Low-level (e.g., C, C++, Rust). Directly manages system resources.
  - Resource management, performance, efficiency, concurrency.

# User Space vs. Kernel Space


- Kernel space (or kernel mode):
  - Where the Operating System (OS) kernel runs.
  - *Privileged*: Has full, unrestricted access to all hardware
  - If it crashes, the whole system crashes.
- User space (or user mode):
  - Where your applications run (e.g., bash, Chrome, your C program).
  - *Unprivileged*: Cannot directly access hardware.
  - If it crashes, the OS cleans it up. The system survives.
  - Lives in a "container" of its own virtual memory.

**Q: If user mode is unprivileged, how does it print to the screen?**

# System Calls

## A: Must ask the kernel (OS) to do it.

- A *system call* is a formal request from a user-space program for a kernel-level service.
  - User program hands over a request by making a system call,
  - Kernel goes into the back, does the privileged work, and
  - Returns the result to the user program.



# System calls

- Types of system calls
  - Process management
  - Memory management
  - File I/O
  - Inter-process communication
  - Signal handling
  - ...
- Examples:
  - `open()`: Ask the kernel to open a file.
  - `read()`, `write()`: Ask the kernel to move data.
  - `fork()`: Ask the kernel to create a new process.
  - many more ...

# Linux filesystem

- The filesystem is your interface to
  - physical storage (disks) on your machine
  - storage on other machines
  - output devices
  - etc.
- *Everything* in Linux is a file (programs, text, peripheral devices, terminals, ...)
- Provides a *logical* view of the storage devices

# Working directory


- The current directory in which you are working
- `pwd` command: outputs the absolute path (more on this later) of your working directory
- Unless you specify another directory, commands will assume you want to operate on the working directory

# Home directory

- A special place for each user to store personal files
- When you log in, your working directory will be set to your home directory
- Your home directory is represented by the symbol ~ (tilde)
  - The home directory of “user” is represented by ~user


# Linux file hierarchy

- Directories may contain plain files or other directories
- Leads to a tree structure for the filesystem
- Root directory: /




# Path names

- Separate directories by /
- *Absolute path*
  - start at root and follow the tree
  - e.g. /home/pyoon/foo.txt



# Path names, cont'd

- ***Relative path***
  - start at working directory
  - `..` refers to level above;
  - `.` refers to working dir.
  - If `/home/pyoon/cpsc275` is working dir, all these refer to the same file
    - `../foo.txt`
    - `~/foo.txt`
    - `~pyoon/foo.txt`



# Types of files

- Plain (– in the first bit)
  - Most files
  - Includes binary and text files
- Directory (d)
  - A directory is actually a file
  - Points to another set of files
- many others ...

% ls -al						
total 94						
drwxr-xr-x	2	john	doc	512	Jul 10 22:25	.
drwxr-xr-x	4	bin	bin	1024	Jul 8 11:48	..
-rw-r--r--	1	john	doc	136	Jul 8 14:46	.exrc
-rw-r--r--	1	john	doc	833	Jul 8 14:51	.profile
-rw-rw-rw-	1	john	doc	31273	Jul 10 22:25	ch1
-rw-rw-rw-	1	john	doc	0	Jul 10 21:57	ch2

Diagram illustrating the columns of the ls output:

- type (e.g., drwxr-xr-x)
- access modes (e.g., 2)
- # of links (e.g., 4)
- owner (e.g., john)
- group (e.g., doc)
- size (in bytes) (e.g., 512, 1024, 136, 833, 31273, 0)
- modification date and time (e.g., Jul 10 22:25, Jul 8 11:48, Jul 8 14:46, Jul 8 14:51, Jul 10 22:25, Jul 10 21:57)
- name (e.g., ., .., .exrc, .profile, ch1, ch2)

# File permissions

- Permissions used to allow/disallow access to file/directory contents
  - Read (r), write (w), and execute (x)
  - For owner, group, and world (everyone)
- chmod <mode> <file(s)>  
chmod 700 file.txt (only owner can read, write, and execute)  
chmod g+rwx file.txt

# Basic file I/O

- Processes keep a list of open files
- Files can be opened for reading, writing
- Each file is referenced by a *file descriptor* (integer)
- Three files are opened automatically
  - 0:** Standard Input (STDIN\_FILENO)
  - 1:** Standard Output (STDOUT\_FILENO)
  - 2:** Standard Error (STDERR\_FILENO)

# File I/O System Call: read()

```
nbytes = read(fd, buffer, count)
```

- reads up to count bytes from file and place into buffer
- fd: file descriptor
- buffer: pointer to array
- count: number of bytes to read
- returns number of bytes read or -1 if error

# File I/O System Call: write()

```
nbytes = write(fd, buffer, count)
```

- writes count bytes from buffer to a file
- fd: file descriptor
- buffer: pointer to array
- count: number of bytes to write
- returns number of bytes written or -1 if error

# Example: A Simple Copy Program, Version I

```
#define BUFSIZE 1

void main(void)
{
    char buf[BUFSIZE];

    int n;
    while ((n = read(0, buf, BUFSIZE)) > 0)
        write(1, buf, n);
}
```

# File I/O system call: open ()

```
fd = open(path, flags, mode)
```

- path: **string, absolute or relative path**
- flags:
  - O\_RDONLY - open for reading
  - O\_WRONLY - open for writing
  - O\_RDWR - open for reading and writing
  - O\_CREAT - create the file if it doesn't exist
  - O\_TRUNC - truncate the file if it exists
  - O\_APPEND - only write at the end of the file
- mode: **specify access permissions if using O\_CREAT**

# Access Permission Bits

S\_IRUSR User (owner) can read

S\_IWUSR User (owner) can write

S\_IXUSR User (owner) can execute

S\_IRGRP Group can read

S\_IWGRP Group can write

S\_IXGRP Group can execute

S\_IROTH Others can read

S\_IWOTH Others can write

S\_IXOTH Others can execute

# File I/O system call: close()

```
retval = close(fd)
```

- closes an open file descriptor fd
- returns 0 on success, -1 on error

# Copy Program, Version 2

```
#define BUFSIZE 1
void main(void)
{
    char buf[BUFSIZE];
    int fdr, fdw, n;

    if ((fdr = open("foo.txt", O_RDONLY, 0)) < 0)
        exit(-1);
    if ((fdw = open("bar.txt", O_WRONLY, 0)) < 0)
        exit(-1);
    while ((n = read(fdr, buf, BUFSIZE)) > 0)
        write(fdw, buf, n);
    close(fdr);
    close(fdw);
}
```

# How can we Improve Performance?


- Making a system call is several orders of magnitude more expensive than a function call
- Given what we know, are there interesting things we can do at the application layer to speed things up?



File System  
Layering

# Caching in the application

- Applications can use caching to improve performance just like the kernel
- Most I/O has both
  - Spatial locality
  - Temporal locality
  - An application level cache in the form of the Standard I/O library attempts to take advantage of this



File System  
Layering

# ANSI C Standard Library

- Collection of high-level input and output functions
- Standard libraries promote cross-platform compatibility.
- Functions, types, and macros of the standard library are accessed by the `#include` directive.
- Actual I/O libraries can be linked *statically* or *dynamically*.

# ANSI Standard Libraries

<assert.h>	Diagnostics
<ctype.h>	Character class tests
<errno.h>	Error codes reported by library functions
<float.h>	Implementation-defined floating-point limits
<limits.h>	Implementation-defined limits
<locale.h>	Locale-specific information
<math.h>	Mathematical functions
<setjmp.h>	Non-local jumps
<signal.h>	Signals
<stdarg.h>	Variable argument lists
<stddef.h>	Definitions of general use
<stdio.h>	Input and output
<stdlib.h>	Utility functions
<string.h>	String functions
<time.h>	Time and date functions

# I/O data streams

- I/O is not directly supported by C but by a set of standard library functions defined by the ANSI C standard.
  - The stdio.h header file contains function declarations for I/O and preprocessor macros related to I/O.
  - stdio.h does not contain the source code for I/O library functions!
- All C character based I/O is performed on streams.
  - All I/O streams must be opened and closed
  - A sequence of characters received from the keyboard is an example of a text stream

# File I/O

- General-purpose I/O functions allow us to specify the stream on which they act
- Must declare a pointer to a FILE struct for each physical file we want to manipulate

```
FILE* infile;  
FILE* outfile;
```

- The I/O stream is "opened" and the FILE struct instantiated by the **fopen** function

```
infile = fopen ("myinfile", "r");  
outfile = fopen ("myoutfile", "w");
```

- Returns pointer to file descriptor or NULL if unsuccessful

# Standard I/O Streams

- In standard C there are three streams automatically opened upon program execution:

FILE *stdin	// input stream
FILE *stdout	// output stream
FILE *stderr	// error messages

# I/O from Files

- File operation modes are:
  - "r" for reading
  - "w" for writing (an existing file will lose its contents)
  - "a" for appending
  - "r+" for reading and writing
  - "b" for binary data
- Once a file is opened, it can be read from or written to with functions such as
  - `fgetc` Read character from stream
  - `fputc` Write character to stream
  - `fread` Read block of data from stream
  - `fwrite` Write block of data to stream
  - `fseek` Reposition stream position indicator
  - `fscanf` Read formatted data from stream
  - `fprintf` Write formatted output to stream

# stdio (caching)

- Each Linux I/O call has a corresponding stdio call

`open()` → `fopen()`

`close` → `fclose()`

`read()` → `fread()`

`write()` → `fwrite()`

# Lab 12

- Implement your own `ls -l` (Linux)
- Both require access to the underlying file system.
- Use Linux system calls
  - `opendir()` – open a directory
  - `closedir()` – close a directory
  - `readdir()` – read directory entries

# System Calls vs. Standard Library

- System calls
  - Direct, *unbuffered* requests to the kernel
  - e.g., `write(1, "h", 1)`: 1 syscall per character.
  - Precise control
  - Inefficient (Each syscall has overhead)
- Standard Library functions
  - *Buffered* functions
  - Writes to a user-space buffer.
  - Much more efficient (drastically fewer syscalls).
  - Less precise control (e.g., buffering can be tricky). Must `fflush()` to force buffered output to its destination.

