

Announcements

- Assignment 9
 - Writing a simple shell
 - Due 11:59 p.m. tonight!
- Final Exam
 - 12:00 – 2:00 p.m., Thursday, December 11
 - Comprehensive
 - Lectures 1-28 (70%)
 - Lectures 29-35 (30%)
- Review Session
 - 2:00 – 4:00 p.m., Wednesday, December 10
 - MECC 127

Lecture 35

Floating-Point Representations

CPSC 275
Introduction to Computer Systems

Floating Point

- **Floating point** is the representation used by most computer systems to approximate real numbers.
- Until the mid-80's, companies each developed their own standard for floating point, resulting in a lack of portability for data.
- In 1985, IEEE 754 was published and became the industry standard.
- The floating-point standard can be used to represent values of the form:

$$V = x * 2^y$$

Fractional Decimal

In decimal, we represent a value d as follows:

$$d_m d_{m-1} \dots d_1 d_0 . d_{-1} d_{-2} \dots d_{-n}$$

 *decimal point*

$$d = \sum_{k=-n}^m d_k * 10^k$$

Fractional Binary

Binary works the same way; to represent a value b :

$$b_m b_{m-1} \dots b_1 b_0 . b_{-1} b_{-2} \dots b_{-n}$$

\uparrow *binary point*

$$b = \sum_{k=-n}^m b_k * 2^k$$

Fractional Binary

$$\text{Binary } 101.01 = 1 * 2^2 + 0 * 2^1 + 1 * 2^0 + 0 * 2^{-1} + 1 * 2^{-2} = 5.25$$

In decimal, what happens when we shift the decimal point one place to the right? How about to the left?

It works the same in binary; shift to the right:

$$101.01 \rightarrow 1010.1 = 1 * 2^3 + 1 * 2^1 + 1 * 2^{-1} = 8 + 2 + \frac{1}{2} = 10.5$$

And a shift to the left:

$$101.01 \rightarrow 10.101 = 1 * 2^1 + 1 * 2^{-1} + 1 * 2^{-3} = 2 + \frac{1}{2} + \frac{1}{8} = 2.625$$

Limitations

- In decimal, there are some numbers that cannot be precisely represented in a finite number of decimal places,
 - e.g. $1/3 = 0.333333\dots$
- Likewise in binary. We are limited to values that can be expressed as $x * 2^y$, where x must be able to be represented in a finite number of bits.
- For example, $1/5$ can be represented exactly in decimal as 0.2, but not so in binary.

Representation	Value	Decimal
0.0_2	$\frac{0}{2}$	0.0_{10}
0.01_2	$\frac{1}{4}$	0.25_{10}
0.010_2	$\frac{2}{8}$	0.25_{10}
0.0011_2	$\frac{3}{16}$	0.1875_{10}
0.00110_2	$\frac{6}{32}$	0.1875_{10}
0.001101_2	$\frac{13}{64}$	0.203125_{10}
0.0011010_2	$\frac{26}{128}$	0.203125_{10}
0.00110011_2	$\frac{51}{256}$	0.19921875_{10}

IEEE Floating Point Standard

The standard has three components:

1. A sign bit **S**, where **S** is 0 for (+), 1 for (-).
2. The *mantissa M* (or *significand*),
3. The exponent **E**.

It is interpreted as:

$$V = (-1)^S * M * 2^E$$

IEEE Floating Point Standard

To determine the values of S , M , and E for a decimal number, say 10.5:

1. S is pretty easy.
2. Write the binary version of the magnitude. We saw before that 10.5 is 1010.1 in binary.
3. Move the binary point to the left or right until there is a single 1 to its left.
 - For this example, that gives us 1.0101. The number of positions moved is $E = +3$. (Moving left gives us a positive value of E , moving right gives us a negative value.)
4. The bits to the right of the binary point are the fractional part of M , 0101.

Q: But what happens to the leading 1?

*A *normalized number* is a number where the most significant digit of the mantissa is non-zero.

IEEE Floating Point Standard

We can verify that the values $S = 0, M = 0101, E = 3$ have the value 10.5 by applying:

$$V = (-1)^S * M * 2^E$$

We obtain:

$$\begin{aligned} V &= (-1)^0 * \textcolor{red}{1.0101} * 2^3 = 1 * \left(1 + \frac{1}{2^2} + \frac{1}{2^4}\right) * 2^3 \\ &= \left(1 + \frac{1}{4} + \frac{1}{16}\right) * 8 \\ &= \frac{21}{16} * 8 \\ &= 10.5 \end{aligned}$$

IEEE Floating Point Standard

Single precision

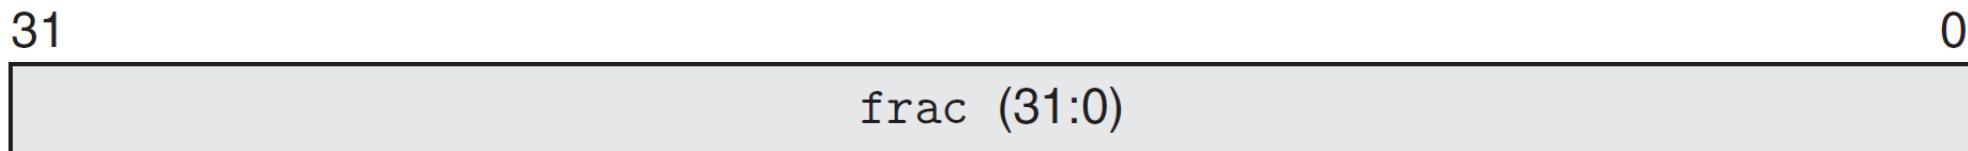


IEEE Floating Point Standard

Single precision



Double precision



Q: But how about negative exponents?

IEEE Floating Point Standard

- To handle negative exponents, a *bias* is introduced.
- If there are k exponent bits, we calculate the bias as $2^{k-1} - 1$.
 - For example, if single precision format has 8 bits for exponent,
$$\text{bias} = 2^7 - 1 = 127$$
- We add the bias to the actual exponent and store the result.
- The stored result is interpreted by subtracting the bias.
- Why introduce bias? (Homework)

Bias for the Exponent

- If we have $k = 4$ bits for the exponent, then the bias will be $2^{k-1} - 1 = 2^3 - 1 = 7$
- To determine (biased) bit representation from true exponent, we add 7 (that will ensure a non-negative value).
- To determine true exponent from (biased) bit representation, we subtract 7 (that will allow negative values).

Binary	E + 7	E
1111	15	8
1110	14	7
1101	13	6
1100	12	5
1011	11	4
1010	10	3
1001	9	2
1000	8	1
0111	7	0
0110	6	-1
0101	5	-2
0100	4	-3
0011	3	-4
0010	2	-5
0001	1	-6
0000	0	-7

IEEE Floating Point Standard – 8-bit model

S (1 bit)	E (4 bits)	M (3 bits)
------------------	-------------------	-------------------

Consider the value 6.5. In binary, that's 110.1

So S = 0, E = 2, and M = 1.101

Our bias is $2^{4-1} - 1 = 2^3 - 1 = 7$, so E = 2 + 7 = 9, or binary 1001.

We drop the leading 1 from M, leaving the fractional part 101.

The result is:

0 1 0 0 1 1 0 1

THANK YOU!

